

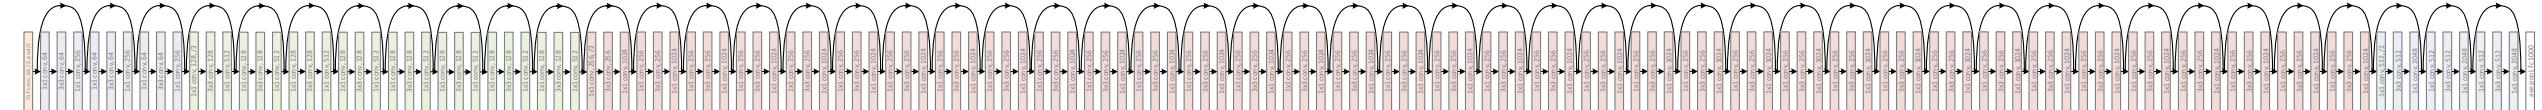
# Deep Learning Gets Way Deeper

## Recent Advances of Deep Learning for Computer Vision

Kaiming He

Research Scientist

Facebook AI Research (FAIR)

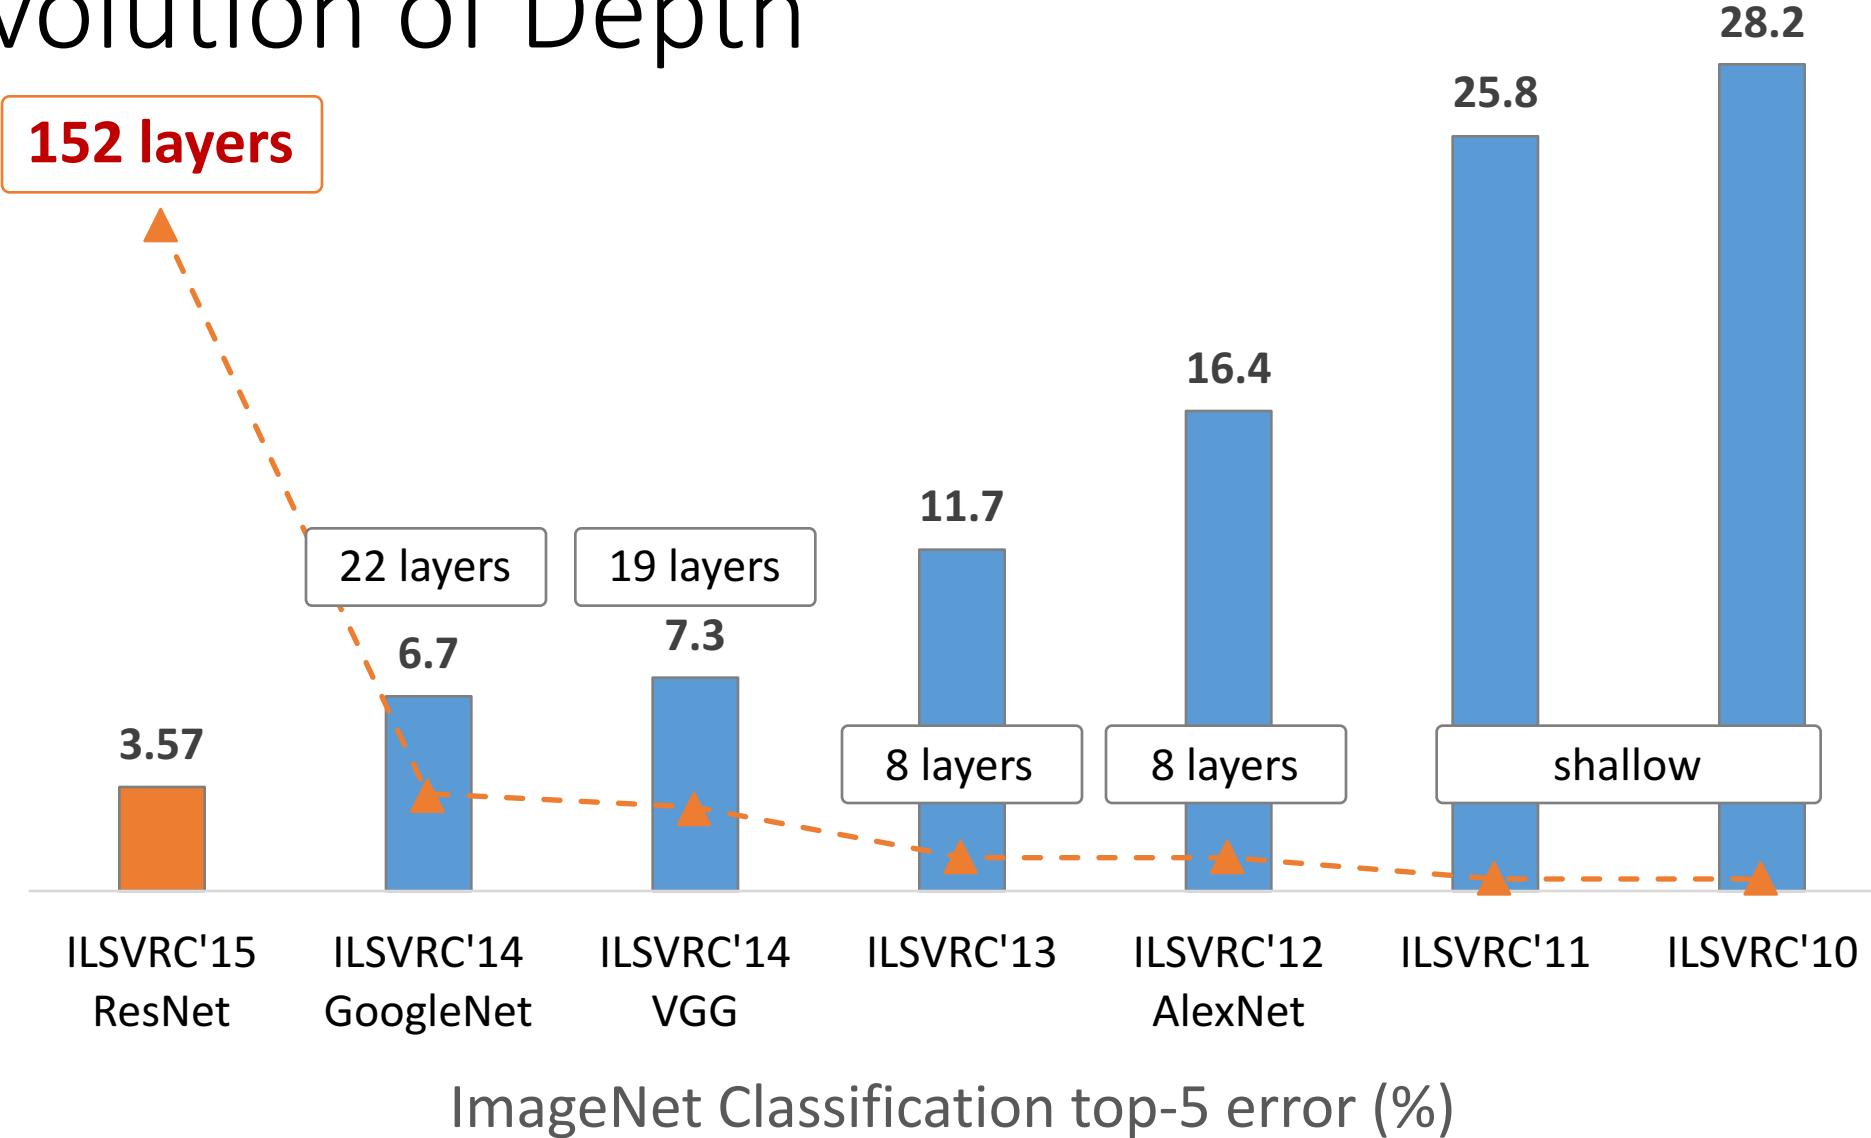


# Overview

- Introduction
  - Look at some recent progress of deep learning for computer vision
- From Shallow Models to 100+ Layers
  - Advances and challenges of getting way deeper
- From Classification to Detection
  - Deep learning for complex recognition applications

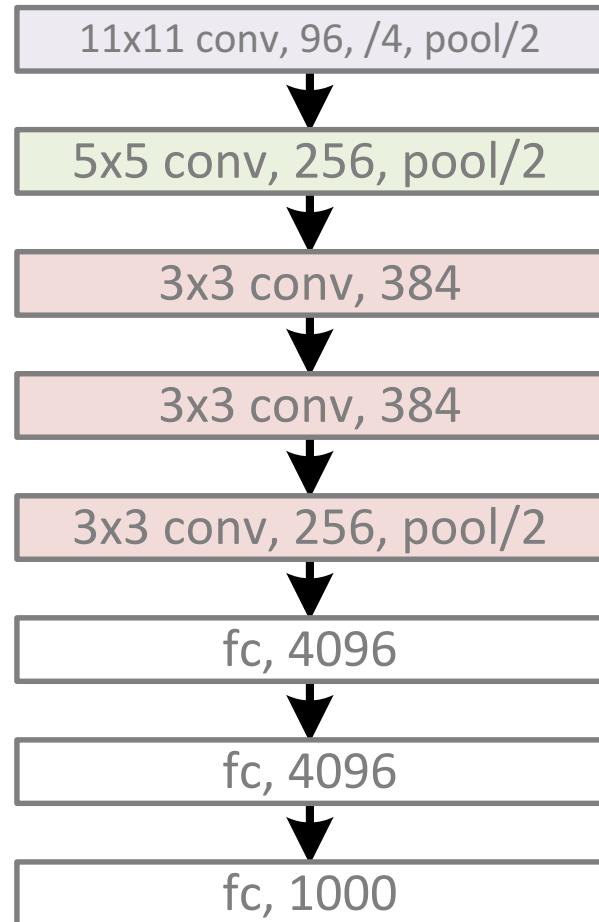
# Introduction

# Revolution of Depth



# Revolution of Depth

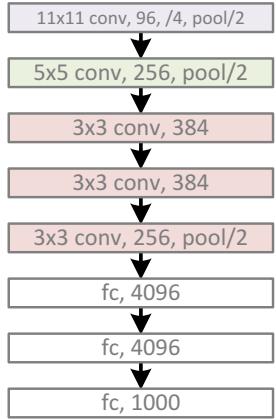
AlexNet, 8 layers  
(ILSVRC 2012)



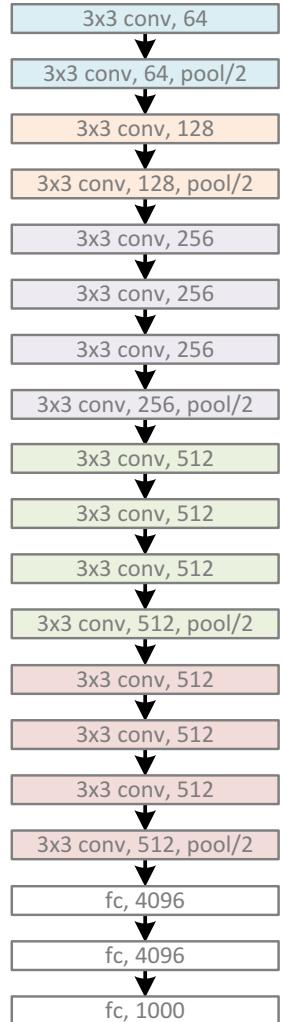
- **ReLU**
- **End-to-end (no pre-training)**
- **Data augmentation**

# Revolution of Depth

AlexNet, 8 layers  
(ILSVRC 2012)



VGG, 19 layers  
(ILSVRC 2014)



- **Very deep**
- **Simply deep**

GoogleNet, 22 layers  
(ILSVRC 2014)



- **Branching**
- **Bottleneck**
- **Skip connection**

# Revolution of Depth

AlexNet, 8 layers  
(ILSVRC 2012)



VGG, 19 layers  
(ILSVRC 2014)

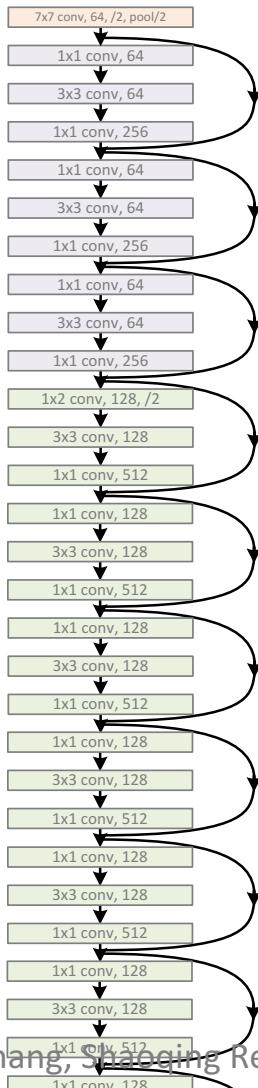


ResNet, **152 layers**  
(ILSVRC 2015)

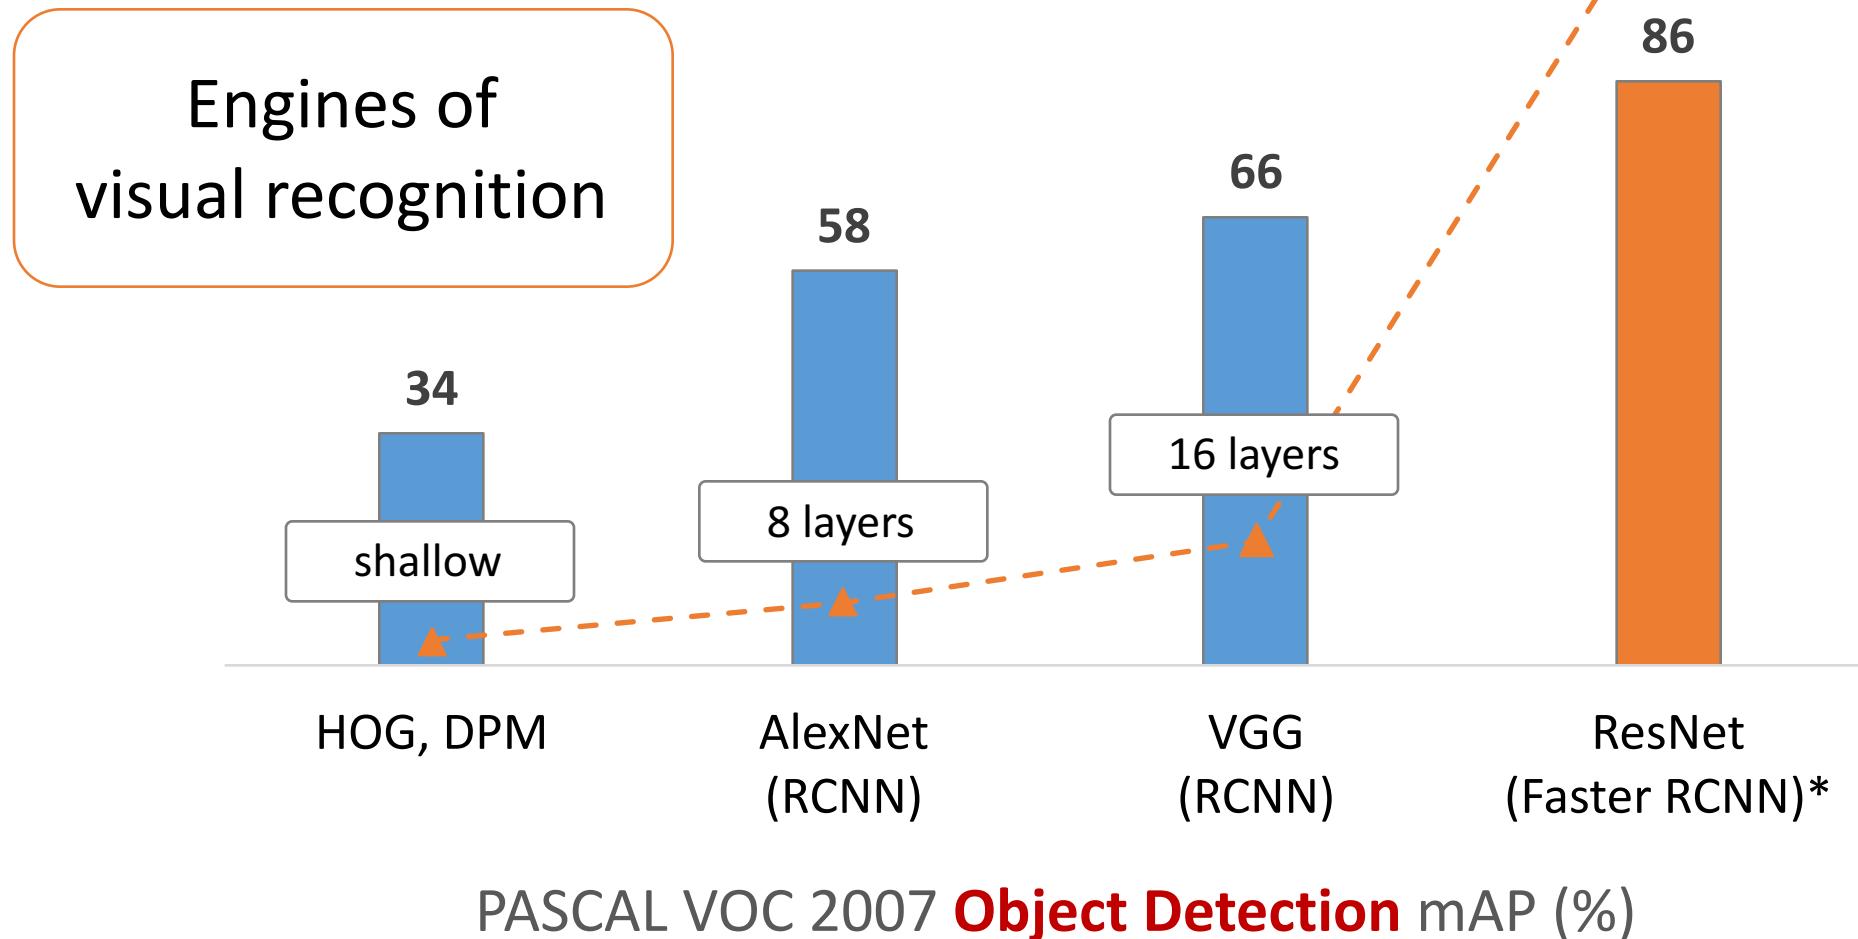


# Revolution of Depth

ResNet, 152 layers



# Revolution of Depth

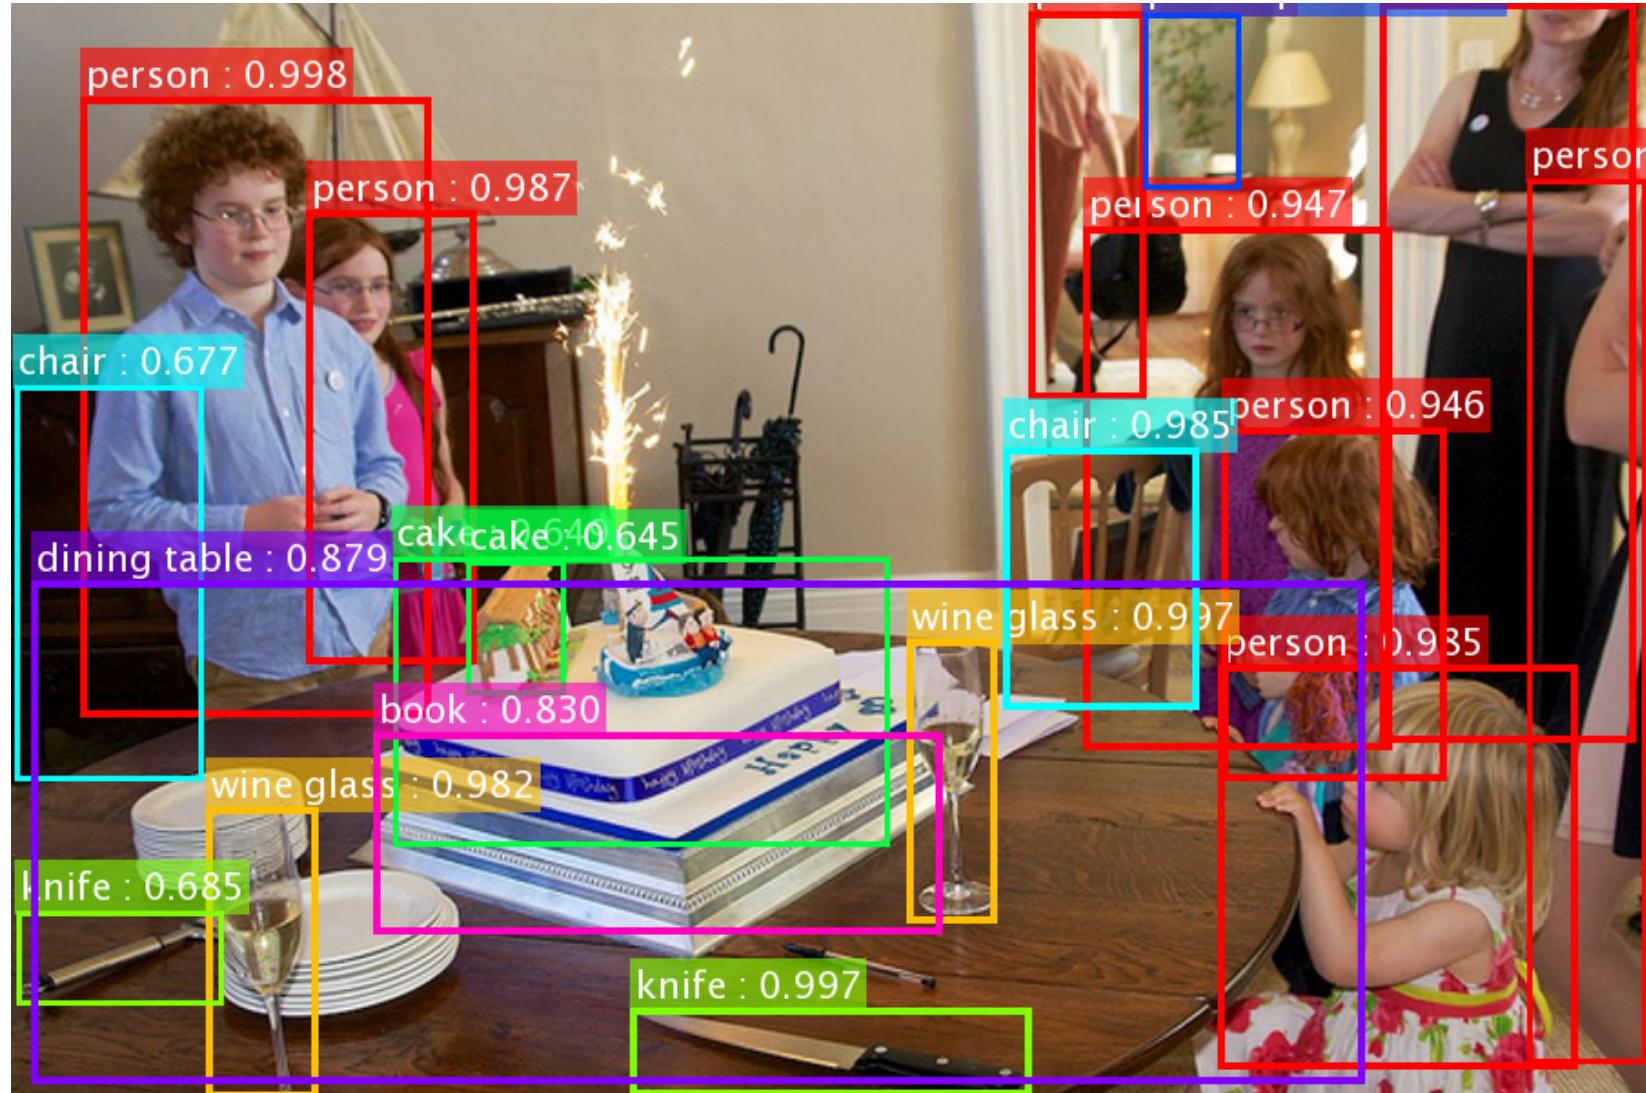


\*w/ other improvements & more data

# ResNets @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**
  - ImageNet Classification: “Ultra-deep” **152-layer** nets
  - ImageNet Detection: **16%** better than 2nd
  - ImageNet Localization: **27%** better than 2nd
  - COCO Detection: **11%** better than 2nd
  - COCO Segmentation: **12%** better than 2nd

\*improvements are relative numbers



## ResNet's object detection result on COCO

\*the original image is from the COCO dataset

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". CVPR 2016.

# Background

From shallow to deep

# Traditional recognition



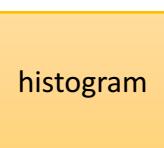
pixels



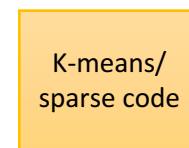
"bus"?



"bus"?



"bus"?



"bus"?

shallow

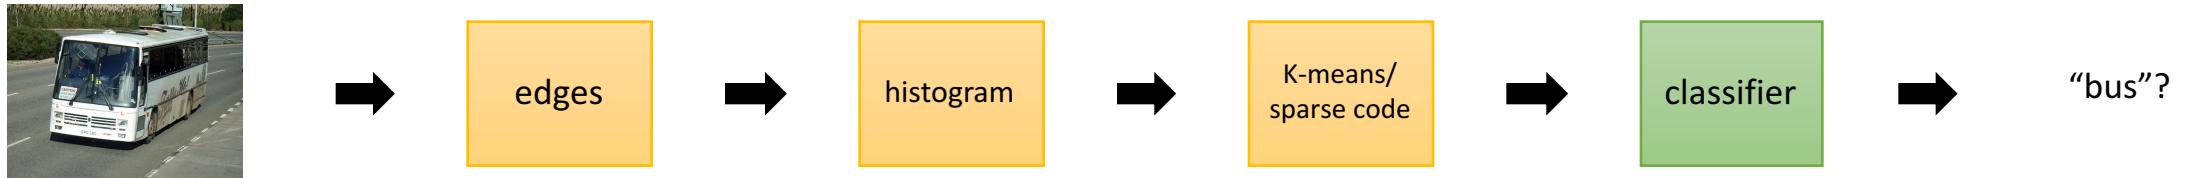
deeper

## But what's next?

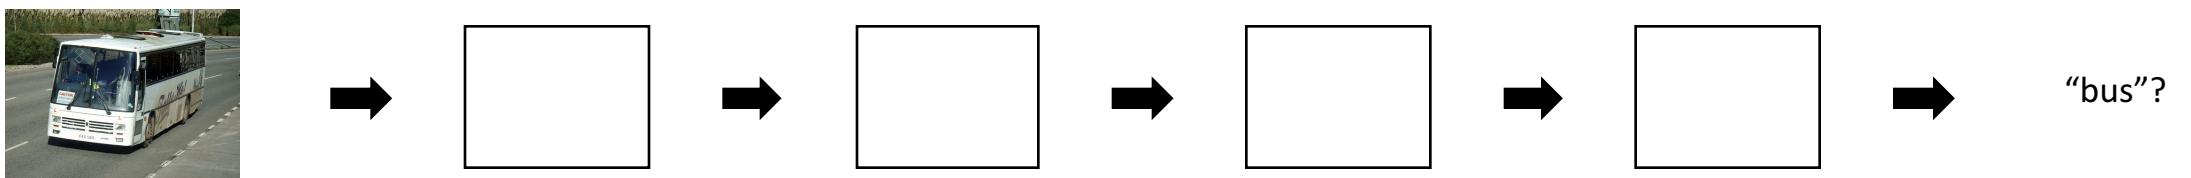


# Deep Learning

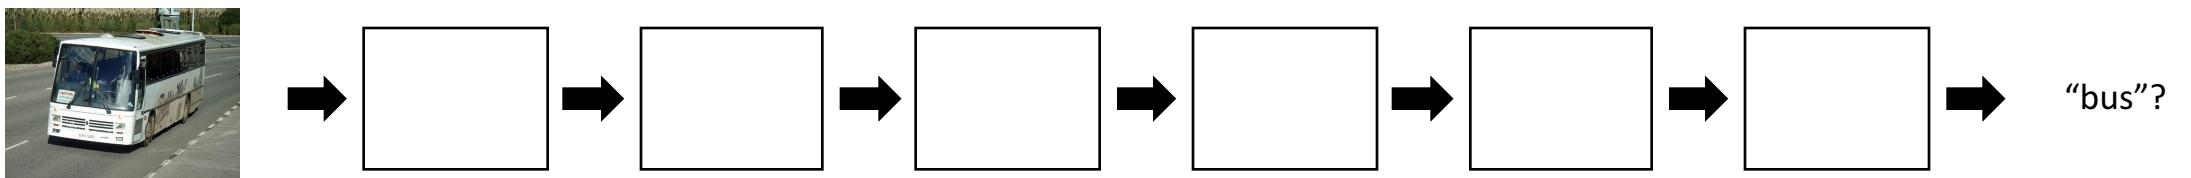
Specialized components



Generic components ("layers")



Repeat elementary layers => Going deeper



- End-to-end learning
- Richer solution space
- **Minimal domain knowledge**

## **Deep Learning is “Easy”**

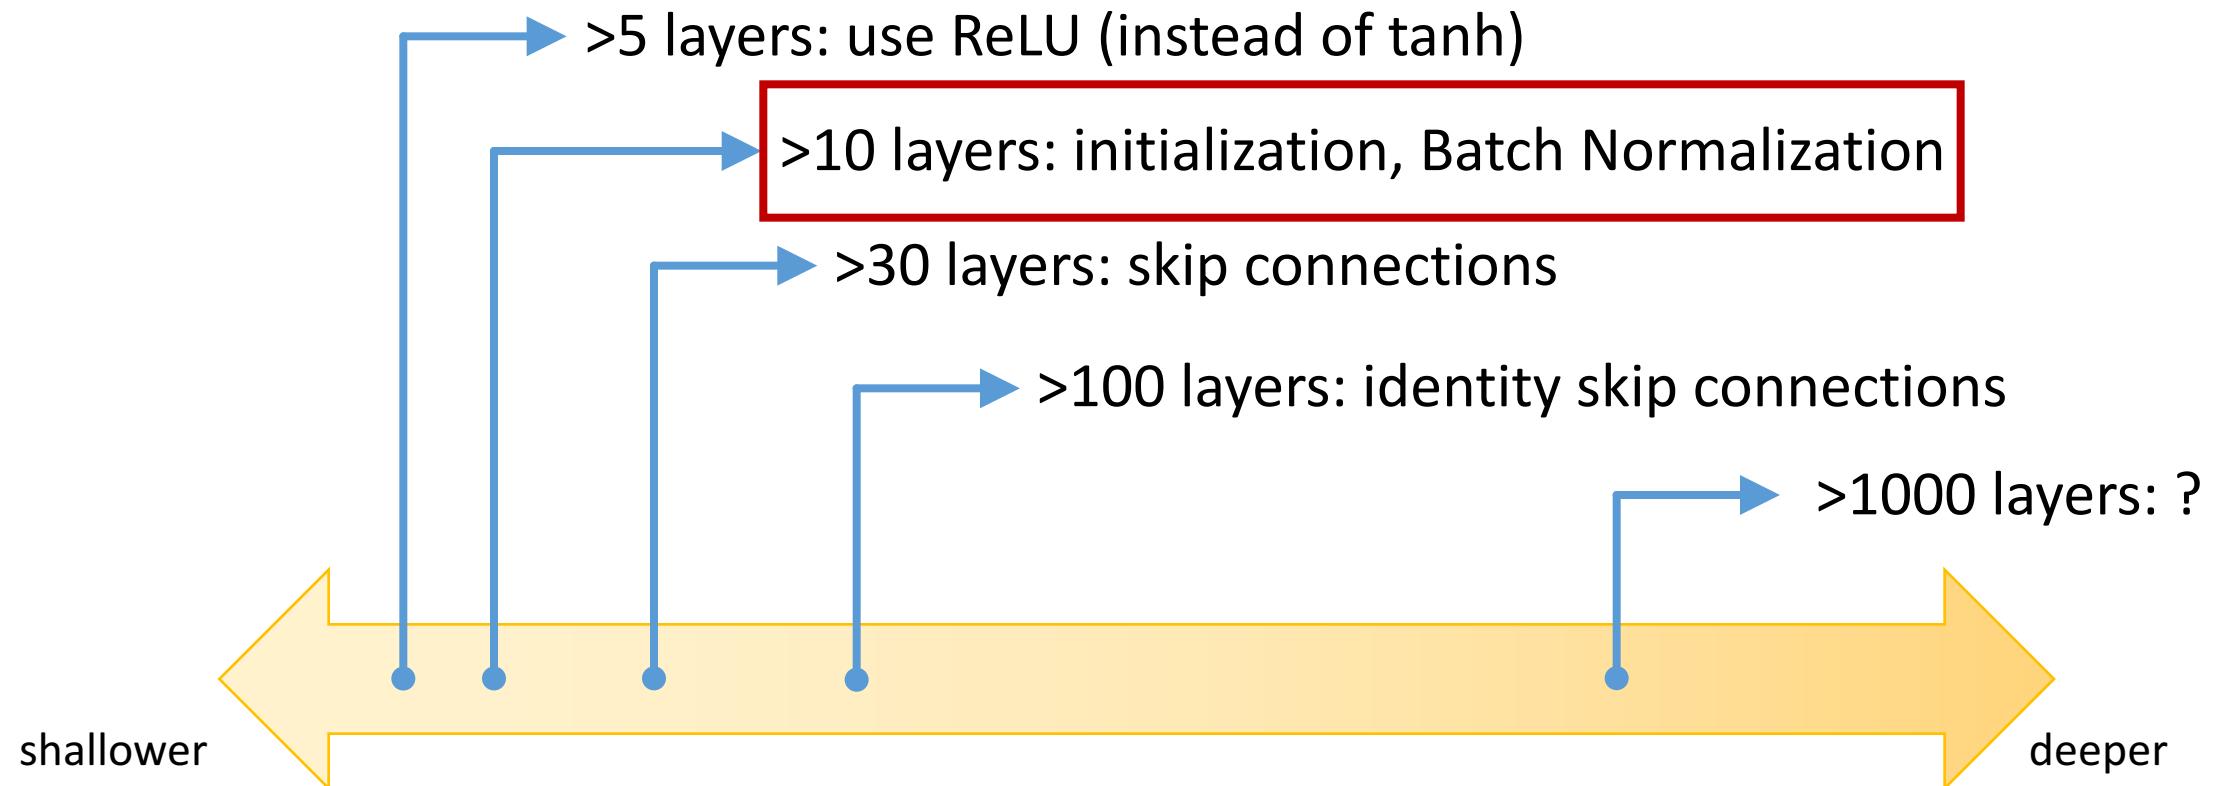
- Minimal domain knowledge
- Data driven
- Features are generalizable

## **Deep Learning is “Hard”**

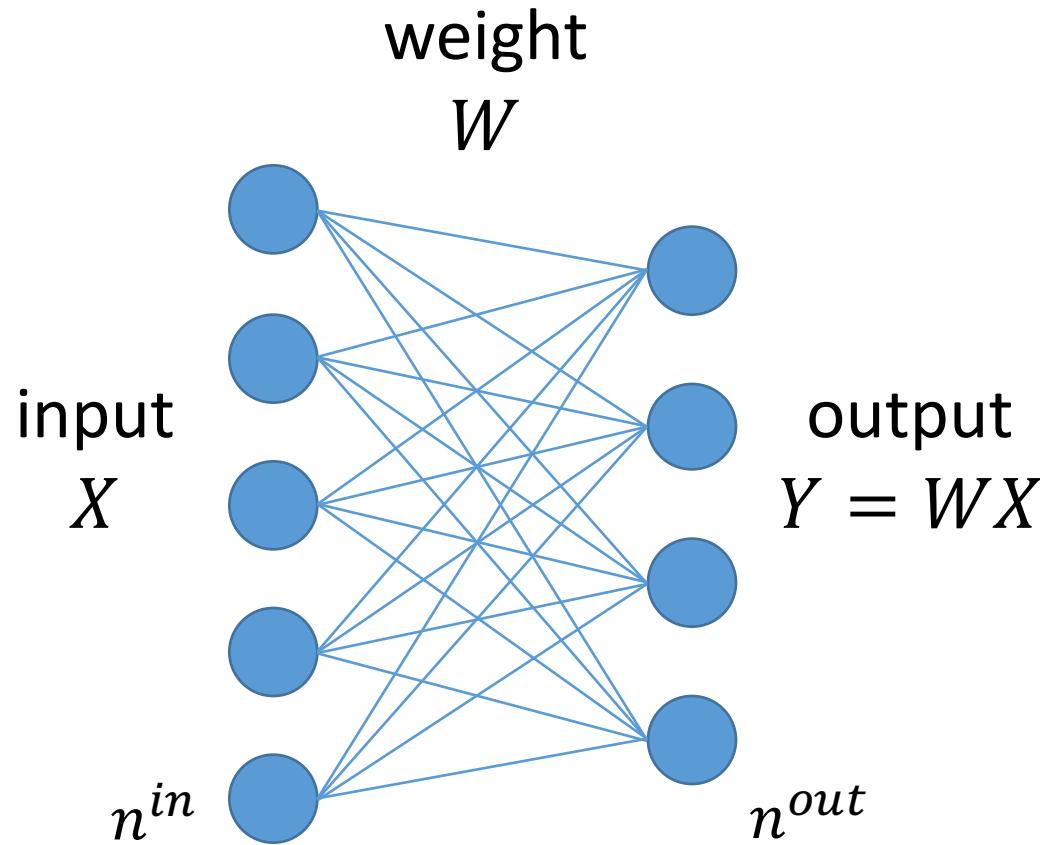
- Black boxes?
- Unstable (vanishing/exploding)?
- Hard to tune hyper-parameters?



# Cheat Sheet of Going Deeper



# Initialization



If:

- Linear activation
- $x, y, w$ : independent

Then:

1-layer:

$$Var[y] = (n^{in} Var[w]) Var[x]$$

Multi-layer:

$$Var[y] = \left( \prod_d n_d^{in} Var[w_d] \right) Var[x]$$

# Initialization

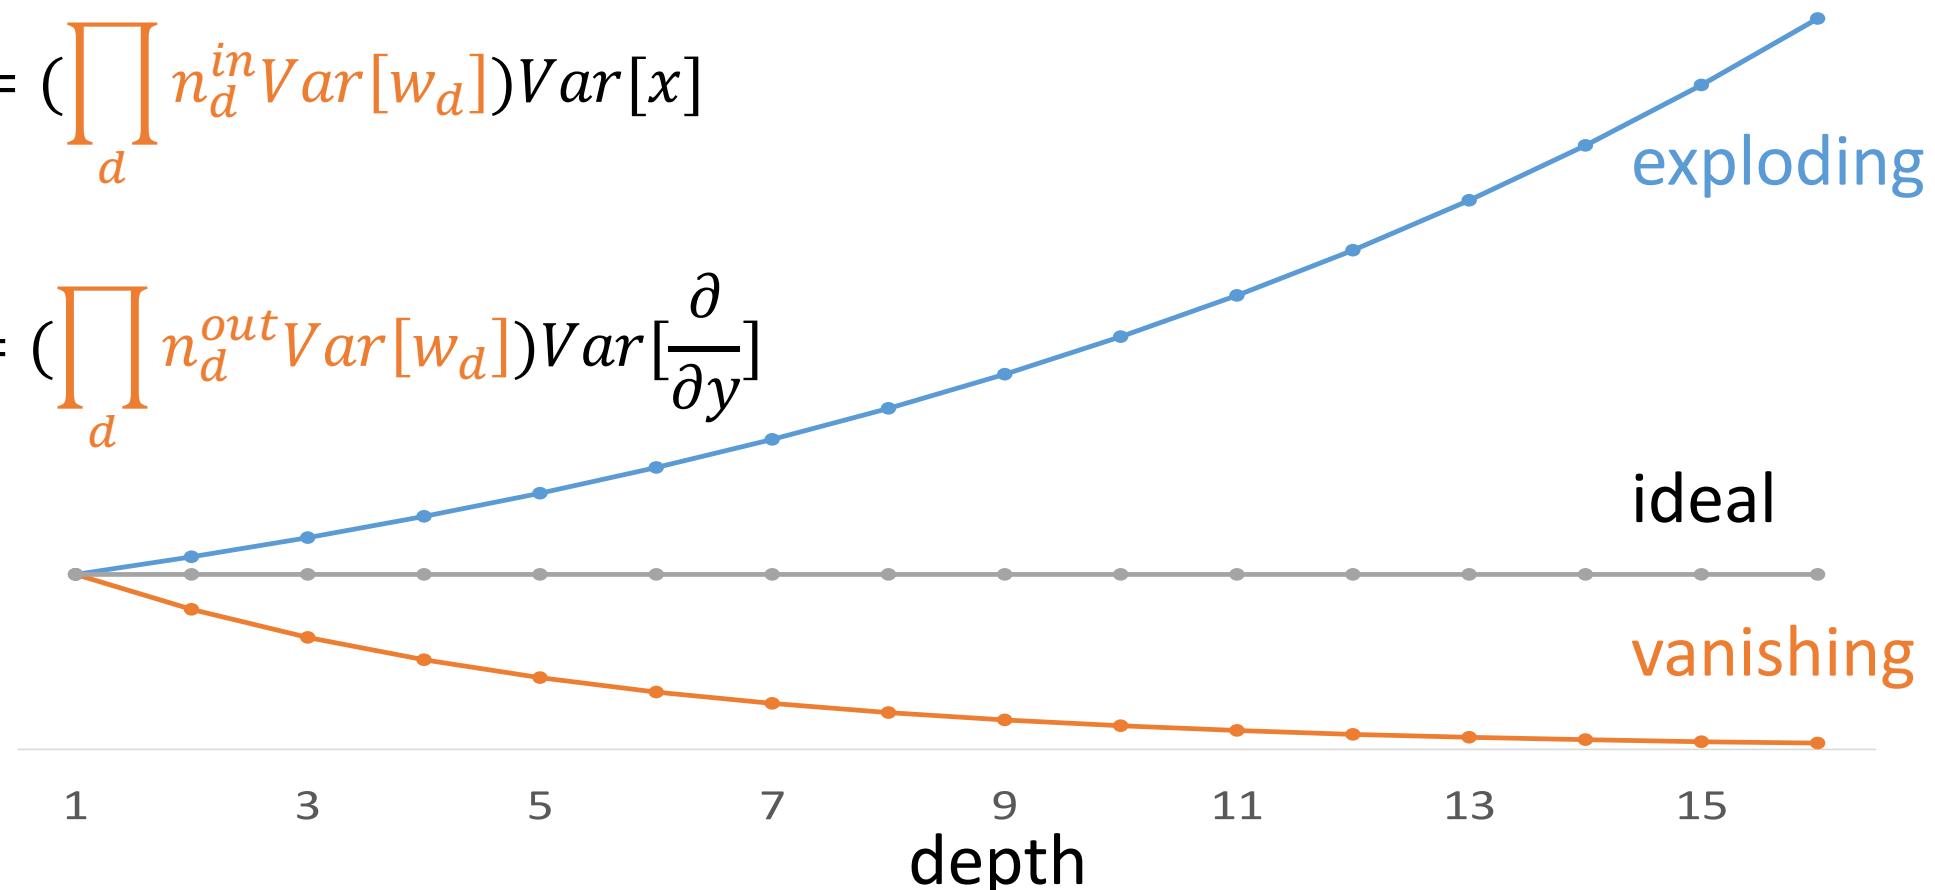
Both forward (response) and backward (gradient) signal can vanish/explode

Forward:

$$Var[y] = \left( \prod_d n_d^{in} Var[w_d] \right) Var[x]$$

Backward:

$$Var\left[\frac{\partial}{\partial x}\right] = \left( \prod_d n_d^{out} Var[w_d] \right) Var\left[\frac{\partial}{\partial y}\right]$$



LeCun et al 1998 "Efficient Backprop"

Glorot & Bengio 2010 "Understanding the difficulty of training deep feedforward neural networks"

# Initialization

- Initialization under **linear** assumption

$$\prod_d n_d^{in} Var[w_d] = const_{fw} \text{ (healthy forward)}$$

and

$$\prod_d n_d^{out} Var[w_d] = const_{bw} \text{ (healthy backward)}$$



$$n_d^{in} Var[w_d] = 1$$

or\*

$$n_d^{out} Var[w_d] = 1$$

\*:  $n_d^{out} = n_{d+1}^{in}$ , so  $\frac{const_{bw}}{const_{fw}} = \frac{n_{last}^{out}}{n_{first}^{in}} < \infty$ .

It is sufficient to use either form.

“Xavier” init in Caffe

LeCun et al 1998 “Efficient Backprop”

Glorot & Bengio 2010 “Understanding the difficulty of training deep feedforward neural networks”

# Initialization

- Initialization under **ReLU**

$$\prod_d \frac{1}{2} n_d^{in} Var[w_d] = const_{fw} \text{ (healthy forward)}$$

and

$$\prod_d \frac{1}{2} n_d^{out} Var[w_d] = const_{bw} \text{ (healthy backward)}$$

$$\frac{1}{2} n_d^{in} Var[w_d] = 1$$

or

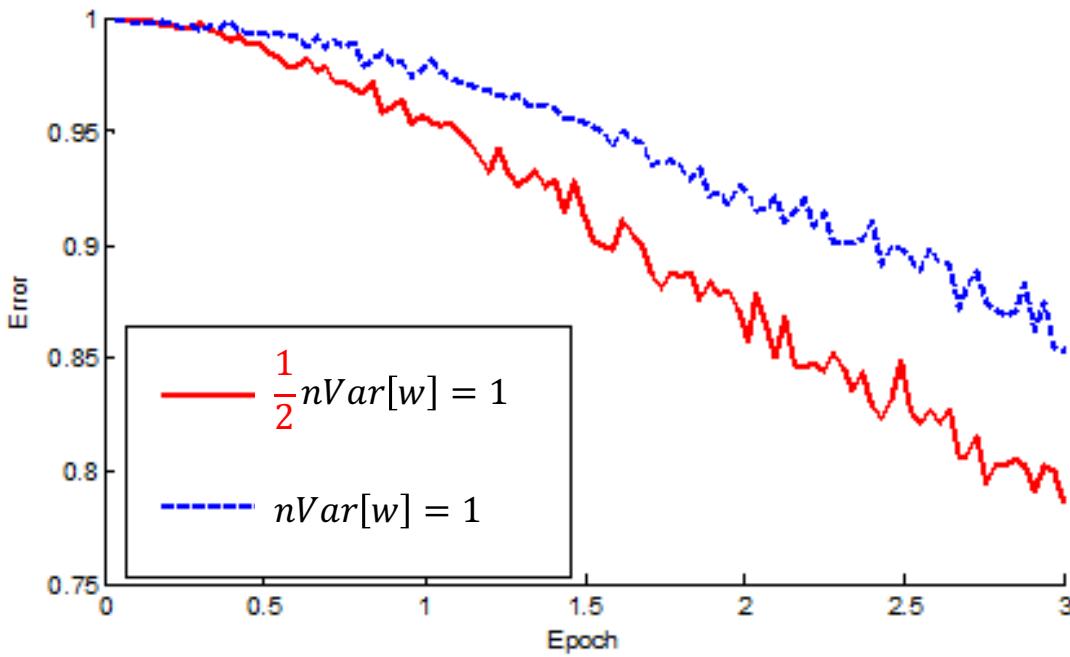
$$\frac{1}{2} n_d^{out} Var[w_d] = 1$$

With  $D$  layers, a factor of 2 per layer has exponential impact of  $2^D$

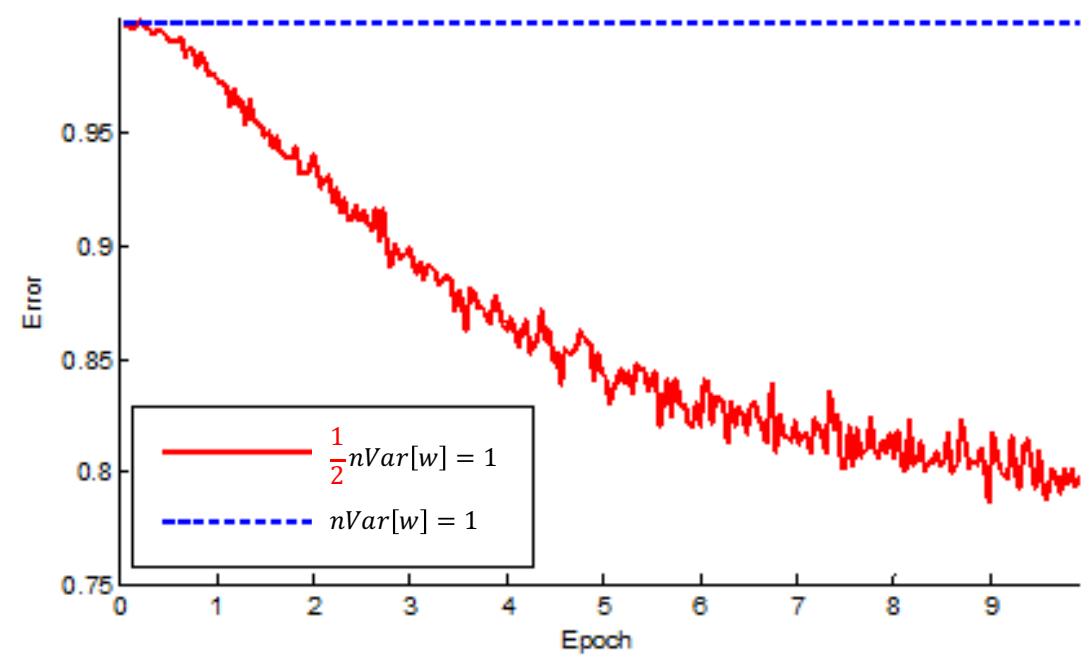
“MSRA” init in Caffe

# Initialization

22-layer ReLU net:  
good init converges faster



30-layer ReLU net:  
good init is able to converge

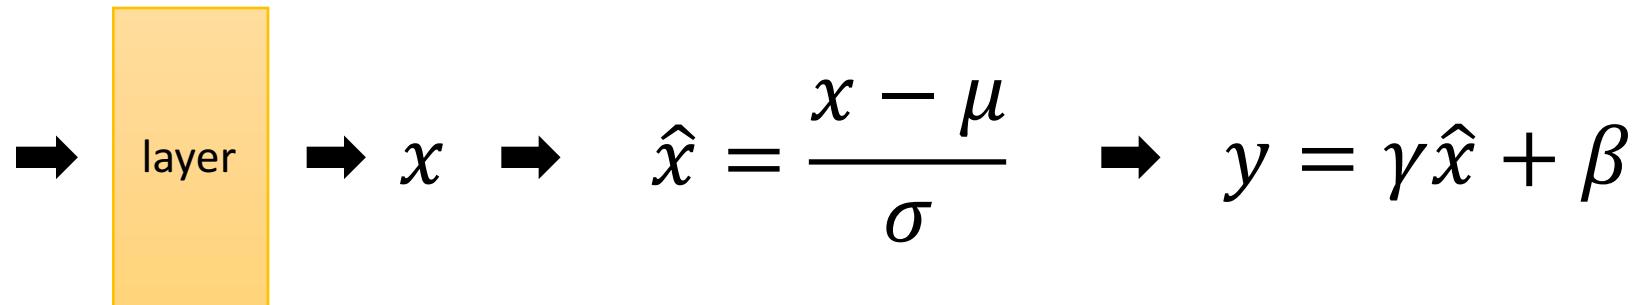


\*Figures show the beginning of training

# Batch Normalization (BN)

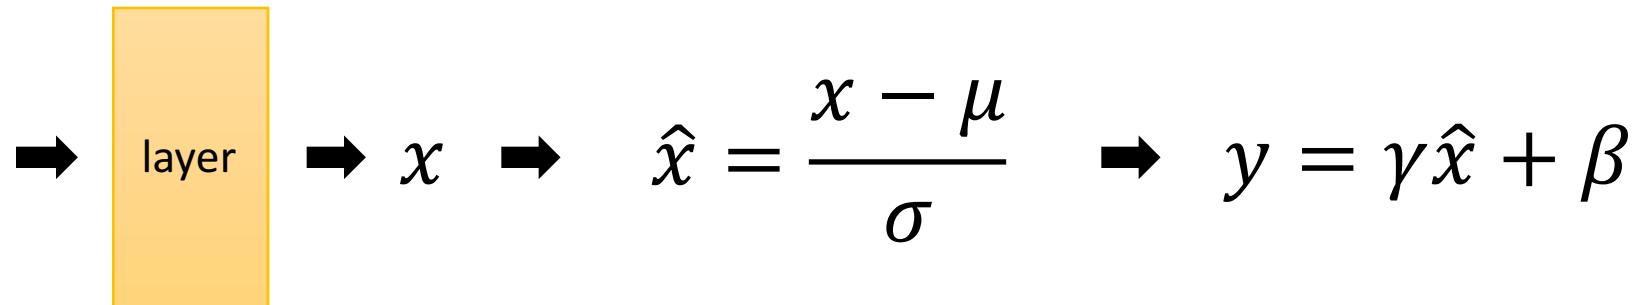
- Normalizing input (LeCun et al 1998 “Efficient Backprop”)
- BN: normalizing **each layer**, for **each mini-batch**
- Greatly accelerate training
- Less sensitive to initialization
- Improve regularization

# Batch Normalization (BN)



- $\mu$ : mean of  $x$  in mini-batch
- $\sigma$ : std of  $x$  in mini-batch
- $\gamma$ : scale
- $\beta$ : shift
- $\mu, \sigma$ : functions of  $x$ ,  
analogous to responses
- $\gamma, \beta$ : parameters to be learned,  
analogous to weights

# Batch Normalization (BN)



2 modes of BN:

- Train mode:
  - $\mu, \sigma$  are functions of  $x$ ; backprop gradients
- Test mode:
  - $\mu, \sigma$  are pre-computed\* on training set

**Caution:** make sure your BN usage is correct

\*: by running average, or post-processing after training

# Batch Normalization (BN)

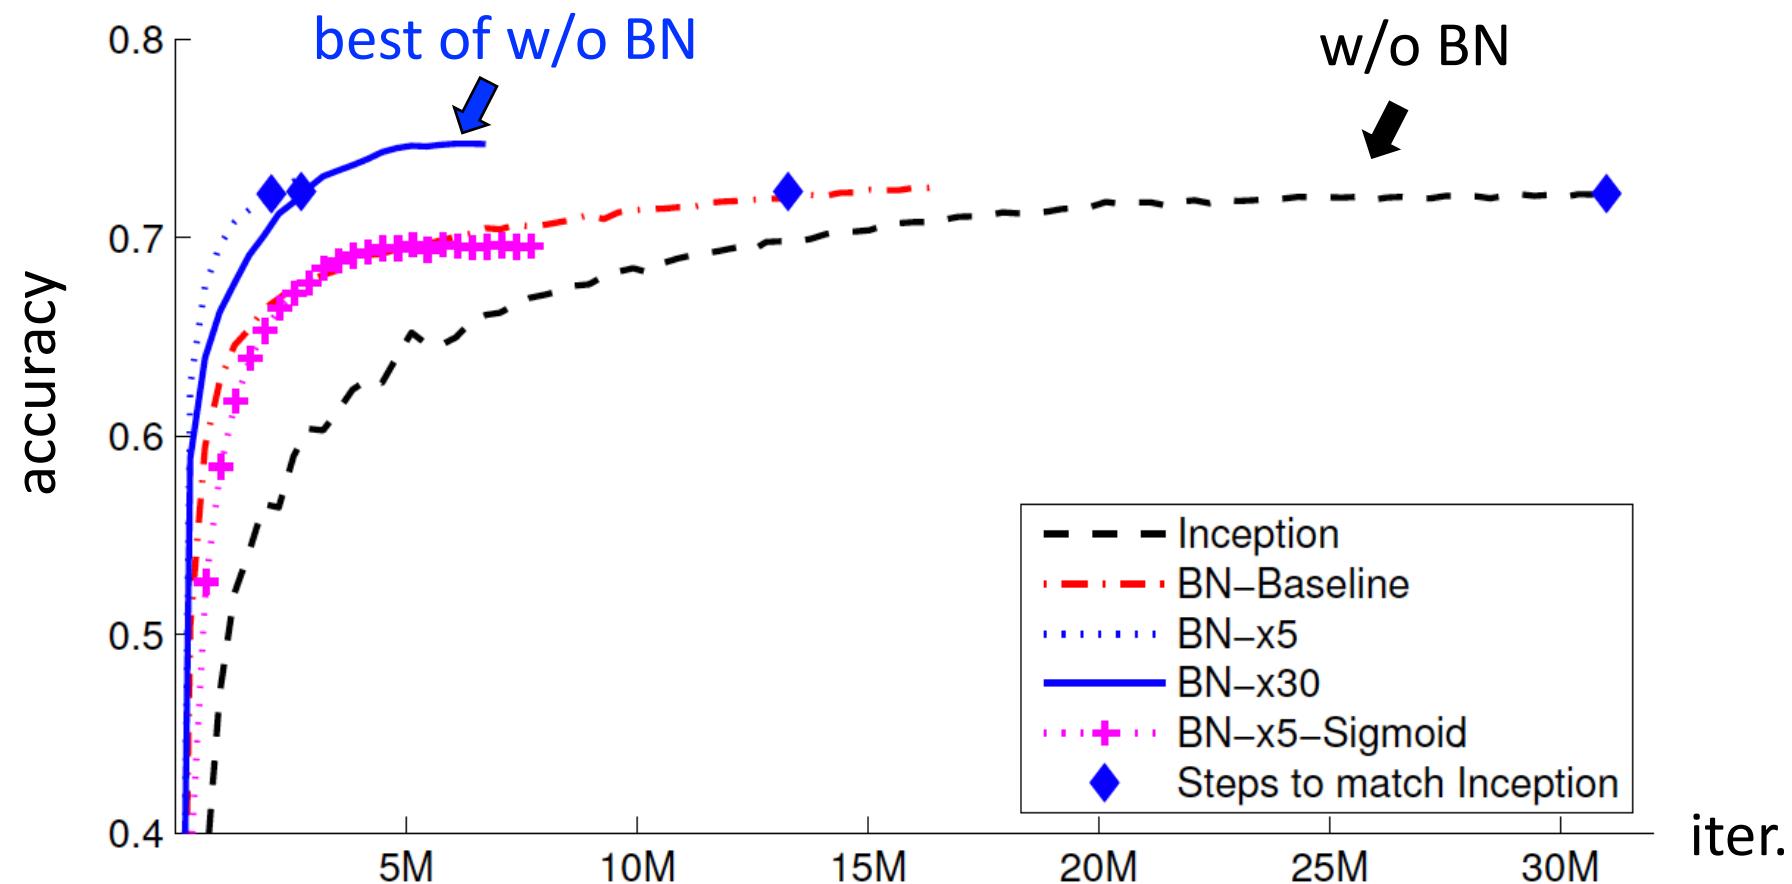


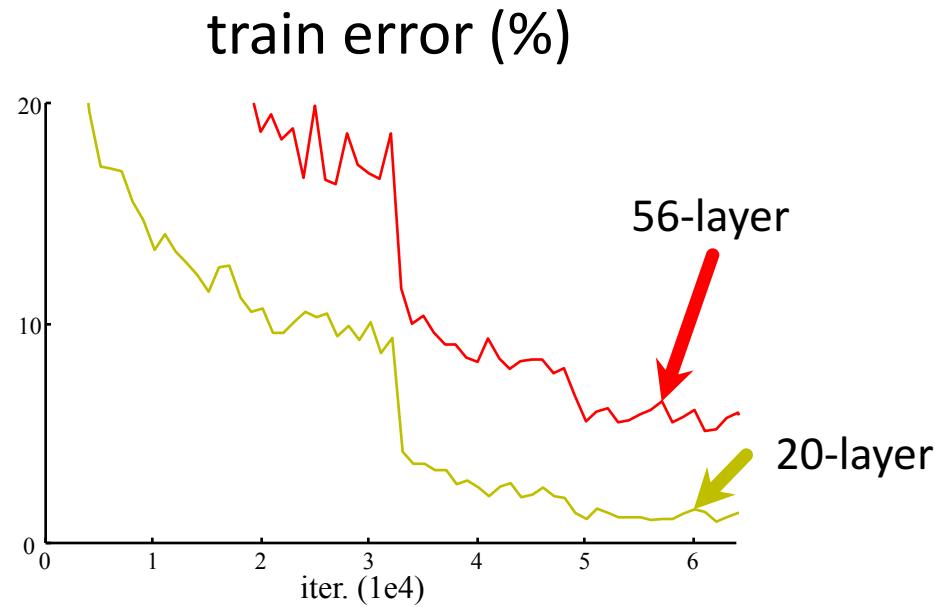
Figure taken from [S. Ioffe & C. Szegedy]

# Deep Residual Networks

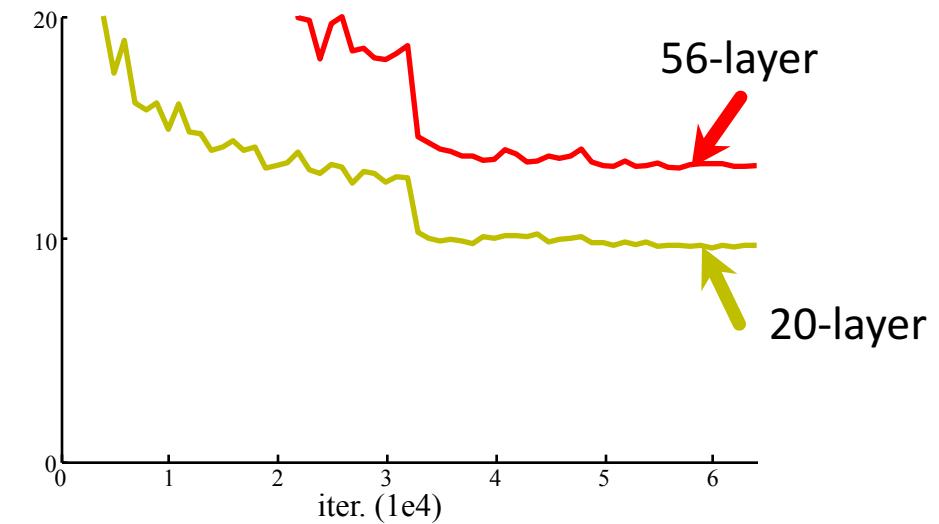
From 10 layers to 100+ layers

# Simply stacking layers?

CIFAR-10

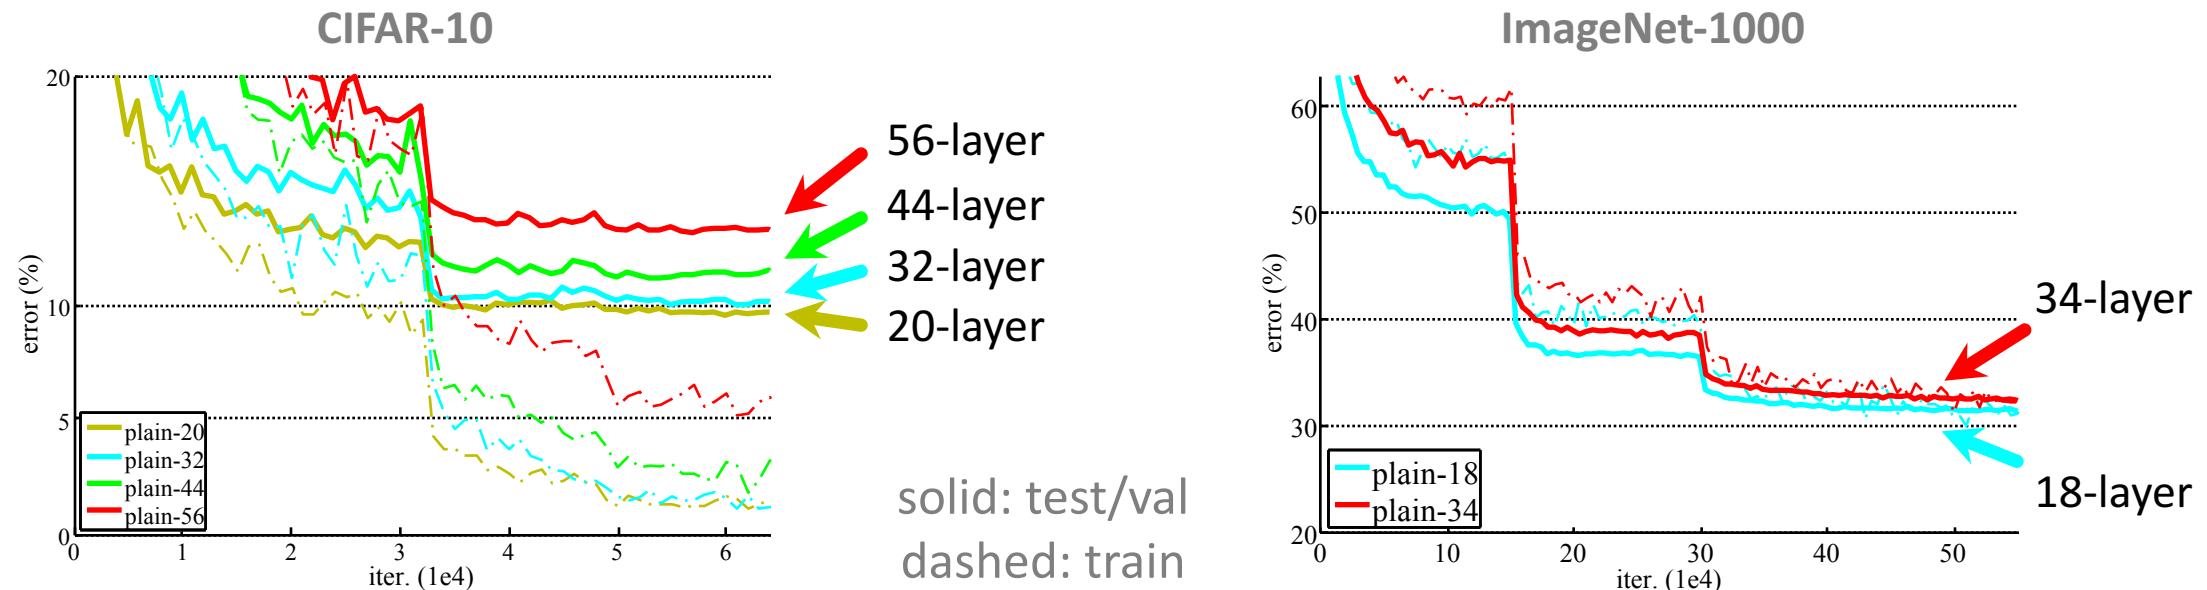


test error (%)



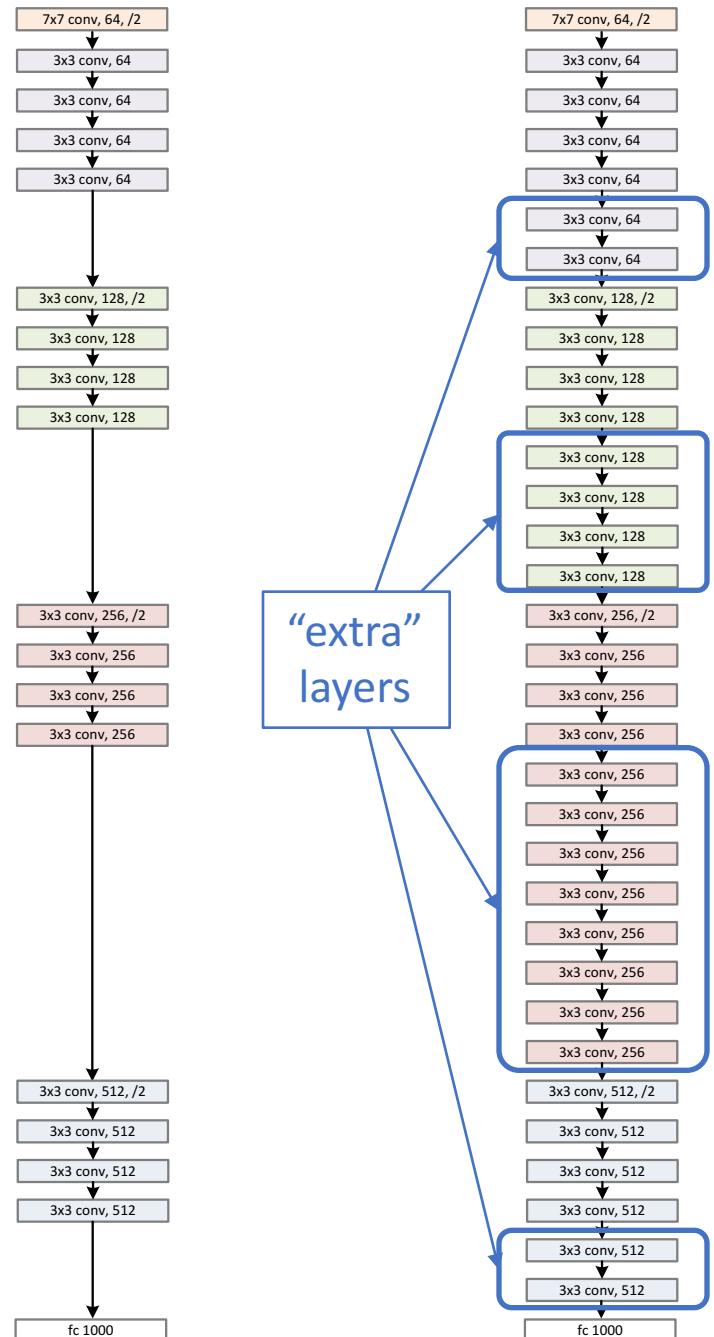
- *Plain* nets: stacking 3x3 conv layers...
- 56-layer net has **higher training error** and test error than 20-layer net

# Simply stacking layers?



- “Overly deep” plain nets have **higher training error**
- A general phenomenon, observed in many datasets

a shallower  
model  
(18 layers)

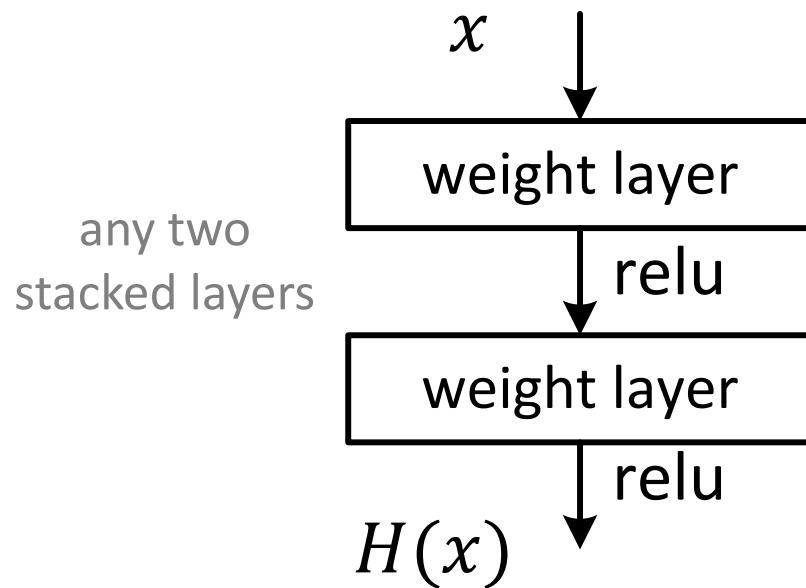


a deeper  
counterpart  
(34 layers)

- Richer solution space
- A deeper model should not have **higher training error**
- A solution *by construction*:
  - original layers: copied from a learned shallower model
  - extra layers: set as **identity**
  - at least the same training error
- **Optimization difficulties**: solvers cannot find the solution when going deeper...

# Deep Residual Learning

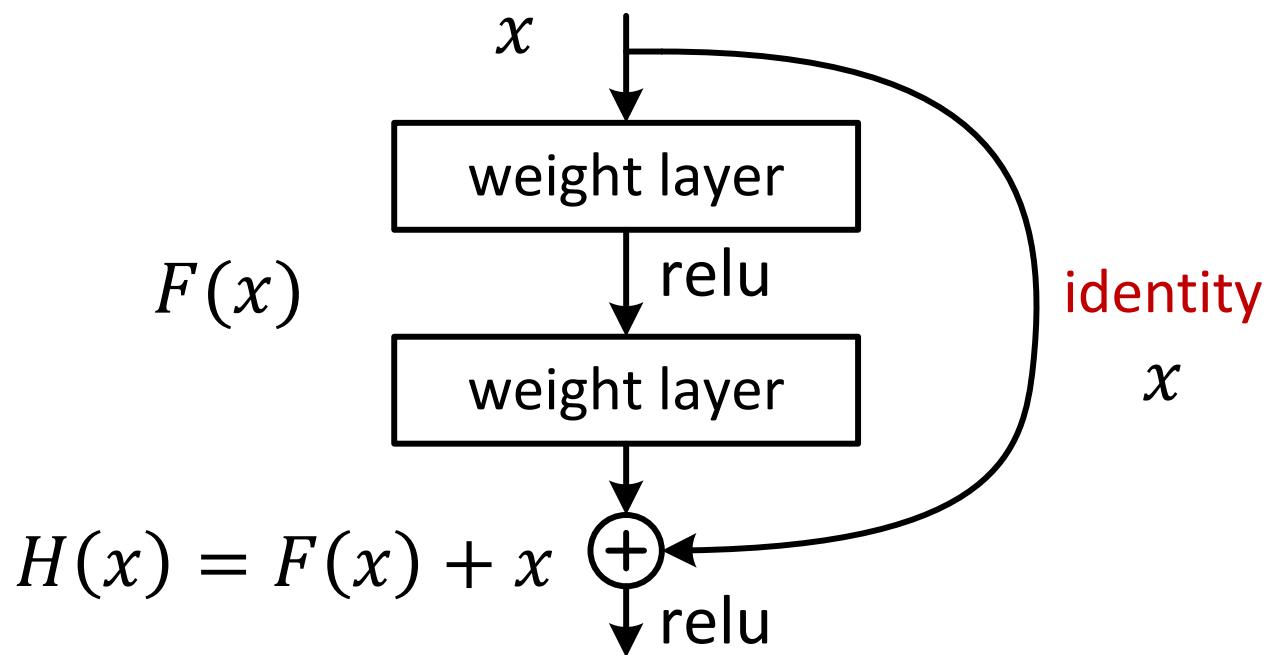
- Plain net



$H(x)$  is any desired mapping,  
hope the 2 weight layers fit  $H(x)$

# Deep Residual Learning

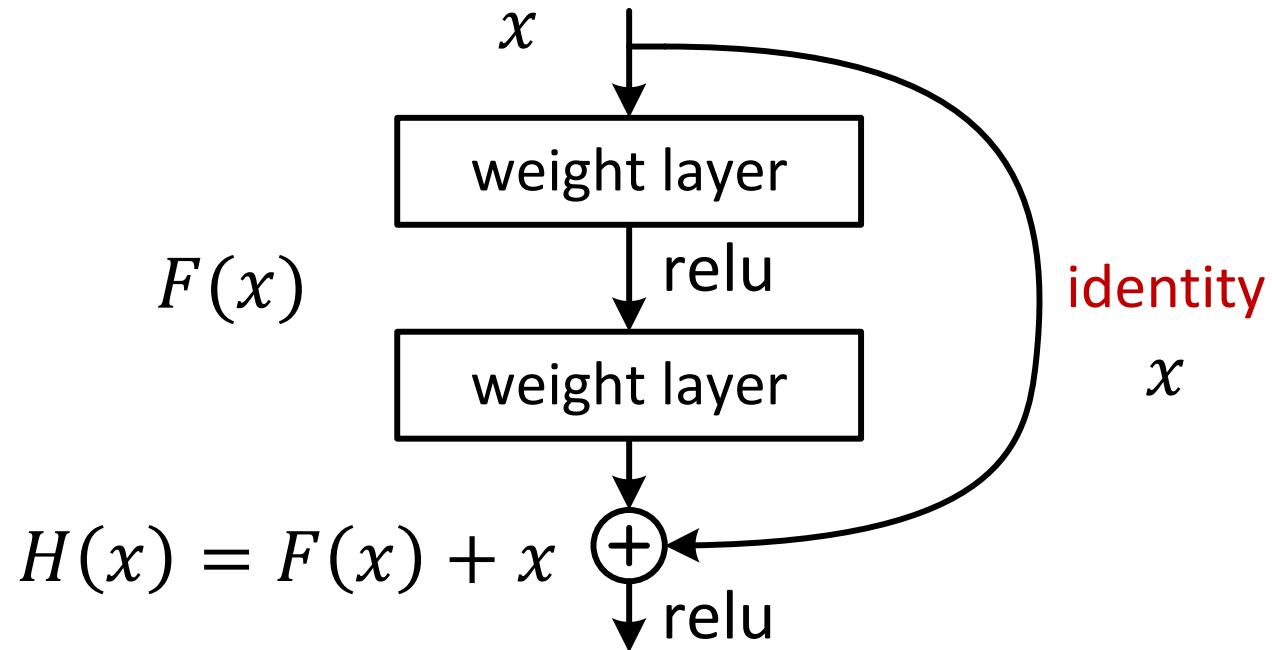
- Residual net



$H(x)$  is any desired mapping,  
hope the 2 weight layers fit  $H(x)$   
hope the 2 weight layers fit  $F(x)$   
let  $H(x) = F(x) + x$

# Deep Residual Learning

- $F(x)$  is a **residual mapping w.r.t. identity**



- If identity were optimal, easy to set weights as 0
- If optimal mapping is closer to identity, easier to find small fluctuations

# Related Works – Residual Representations

- VLAD & Fisher Vector [Jegou et al 2010], [Perronnin et al 2007]
  - Encoding **residual** vectors; powerful shallower representations.
- Product Quantization (IVF-ADC) [Jegou et al 2011]
  - Quantizing **residual** vectors; efficient nearest-neighbor search.
- MultiGrid & Hierarchical Precondition [Briggs, et al 2000], [Szeliski 1990, 2006]
  - Solving **residual** sub-problems; efficient PDE solvers.

# Network “Design”

- Keep it simple

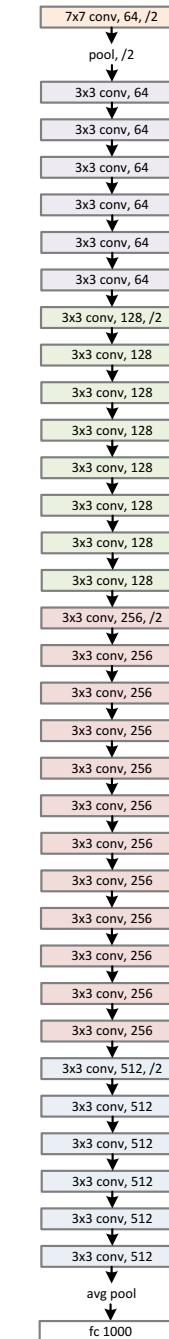
- Our basic design (VGG-style)

- all 3x3 conv (almost)
- spatial size /2 => # filters x2 (~same complexity per layer)
- **Simple design; just deep!**

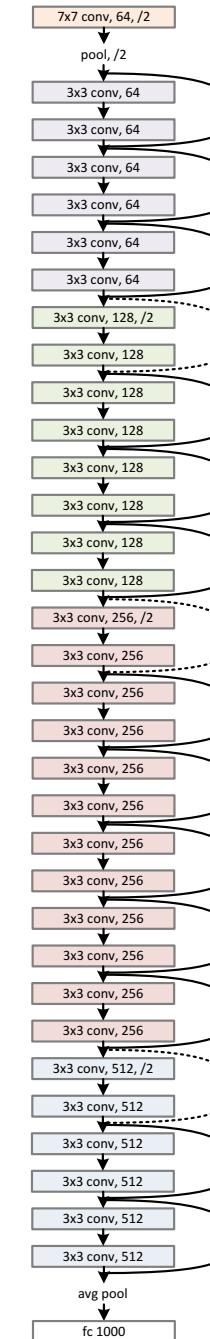
- Other remarks:

- no hidden fc
- no dropout

plain net



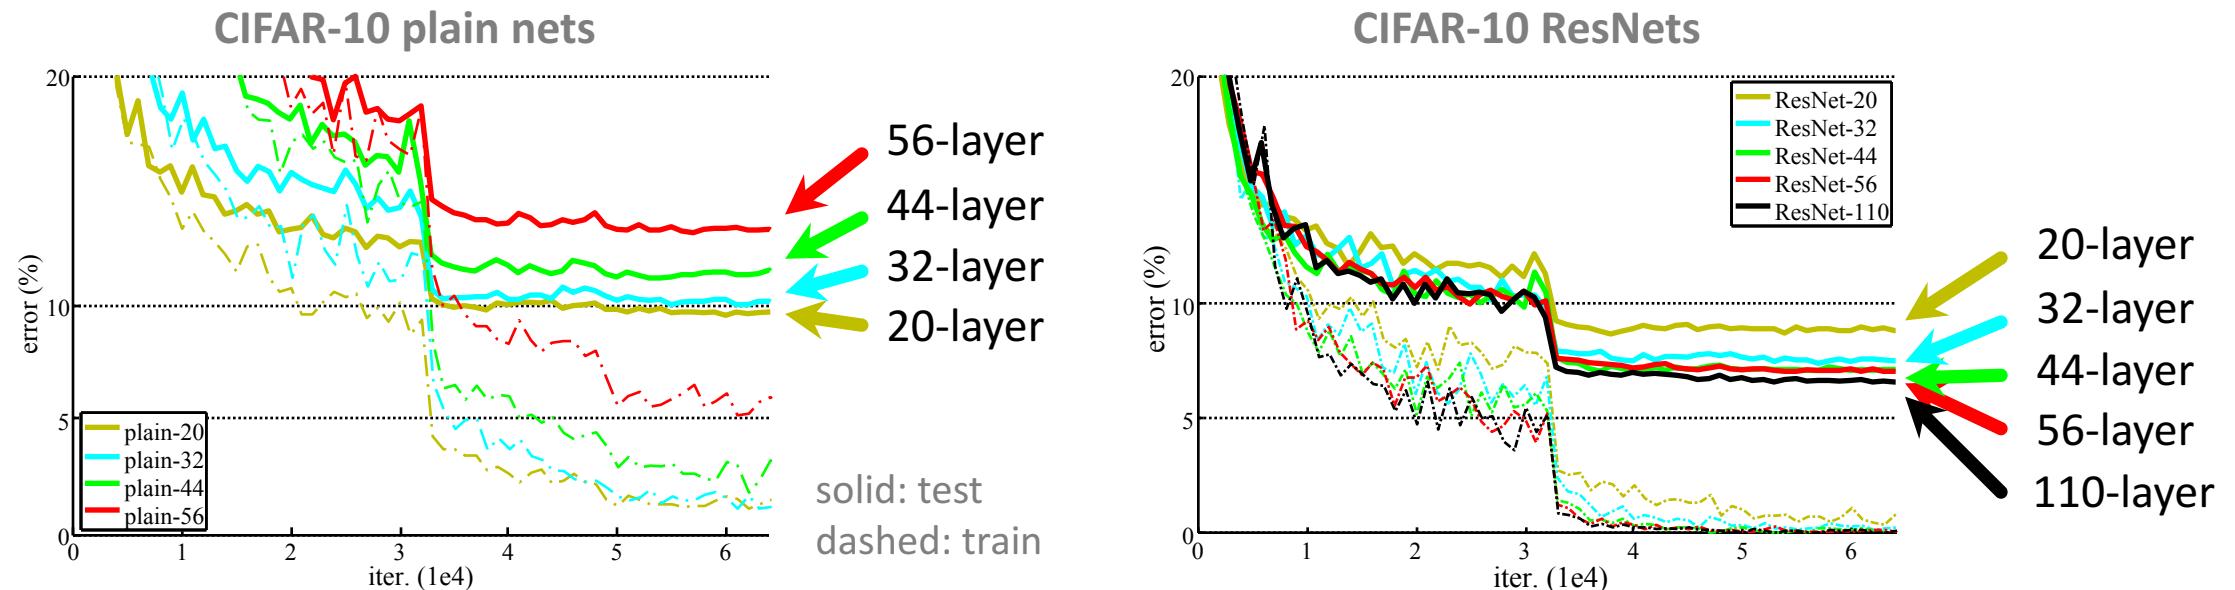
ResNet



# Training

- All plain/residual nets are trained **from scratch**
- All plain/residual nets use Batch Normalization
- Standard hyper-parameters & augmentation

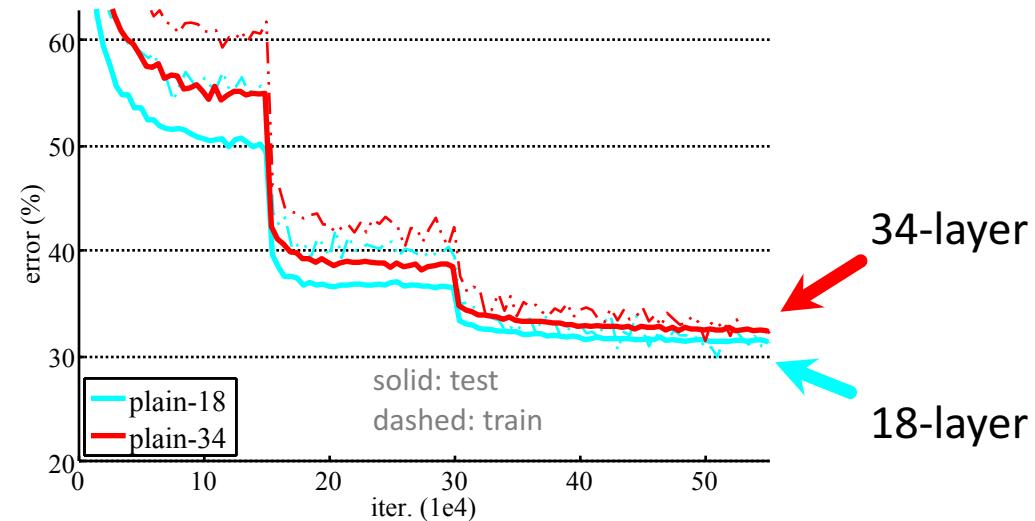
# CIFAR-10 experiments



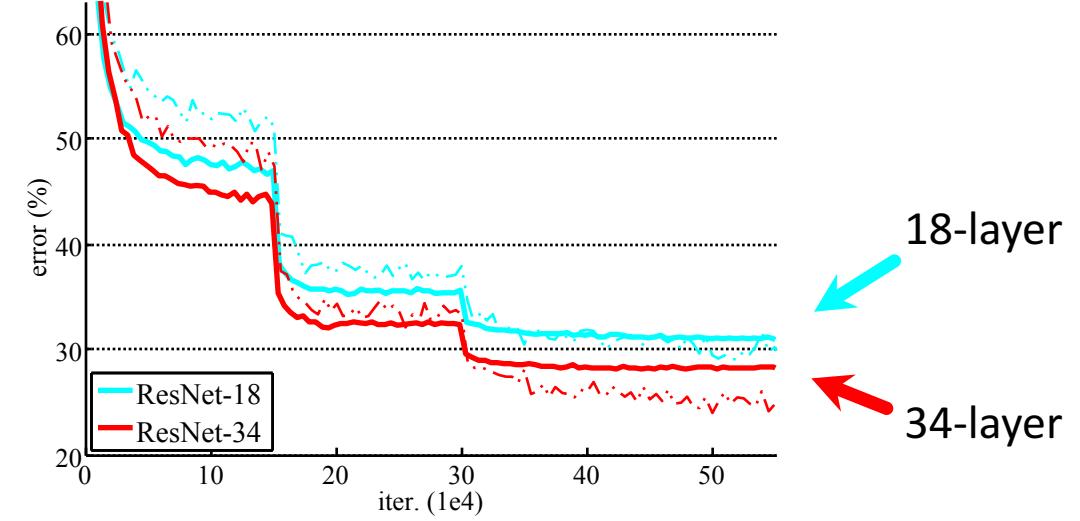
- Deep ResNets can be trained without difficulties
- Deeper ResNets have **lower training error**, and also lower test error

# ImageNet experiments

ImageNet plain nets



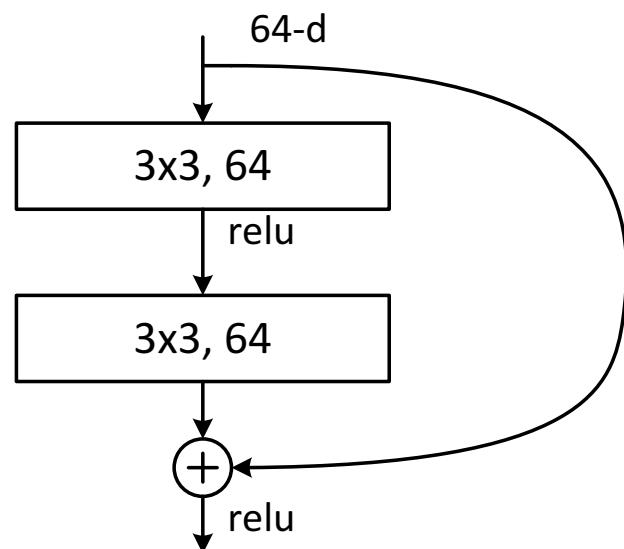
ImageNet ResNets



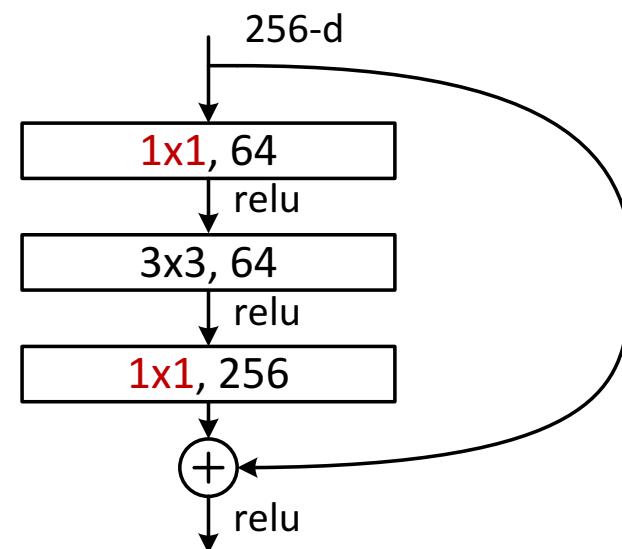
- Deep ResNets can be trained without difficulties
- Deeper ResNets have **lower training error**, and also lower test error

# ImageNet experiments

- A practical design of going deeper



all-3x3

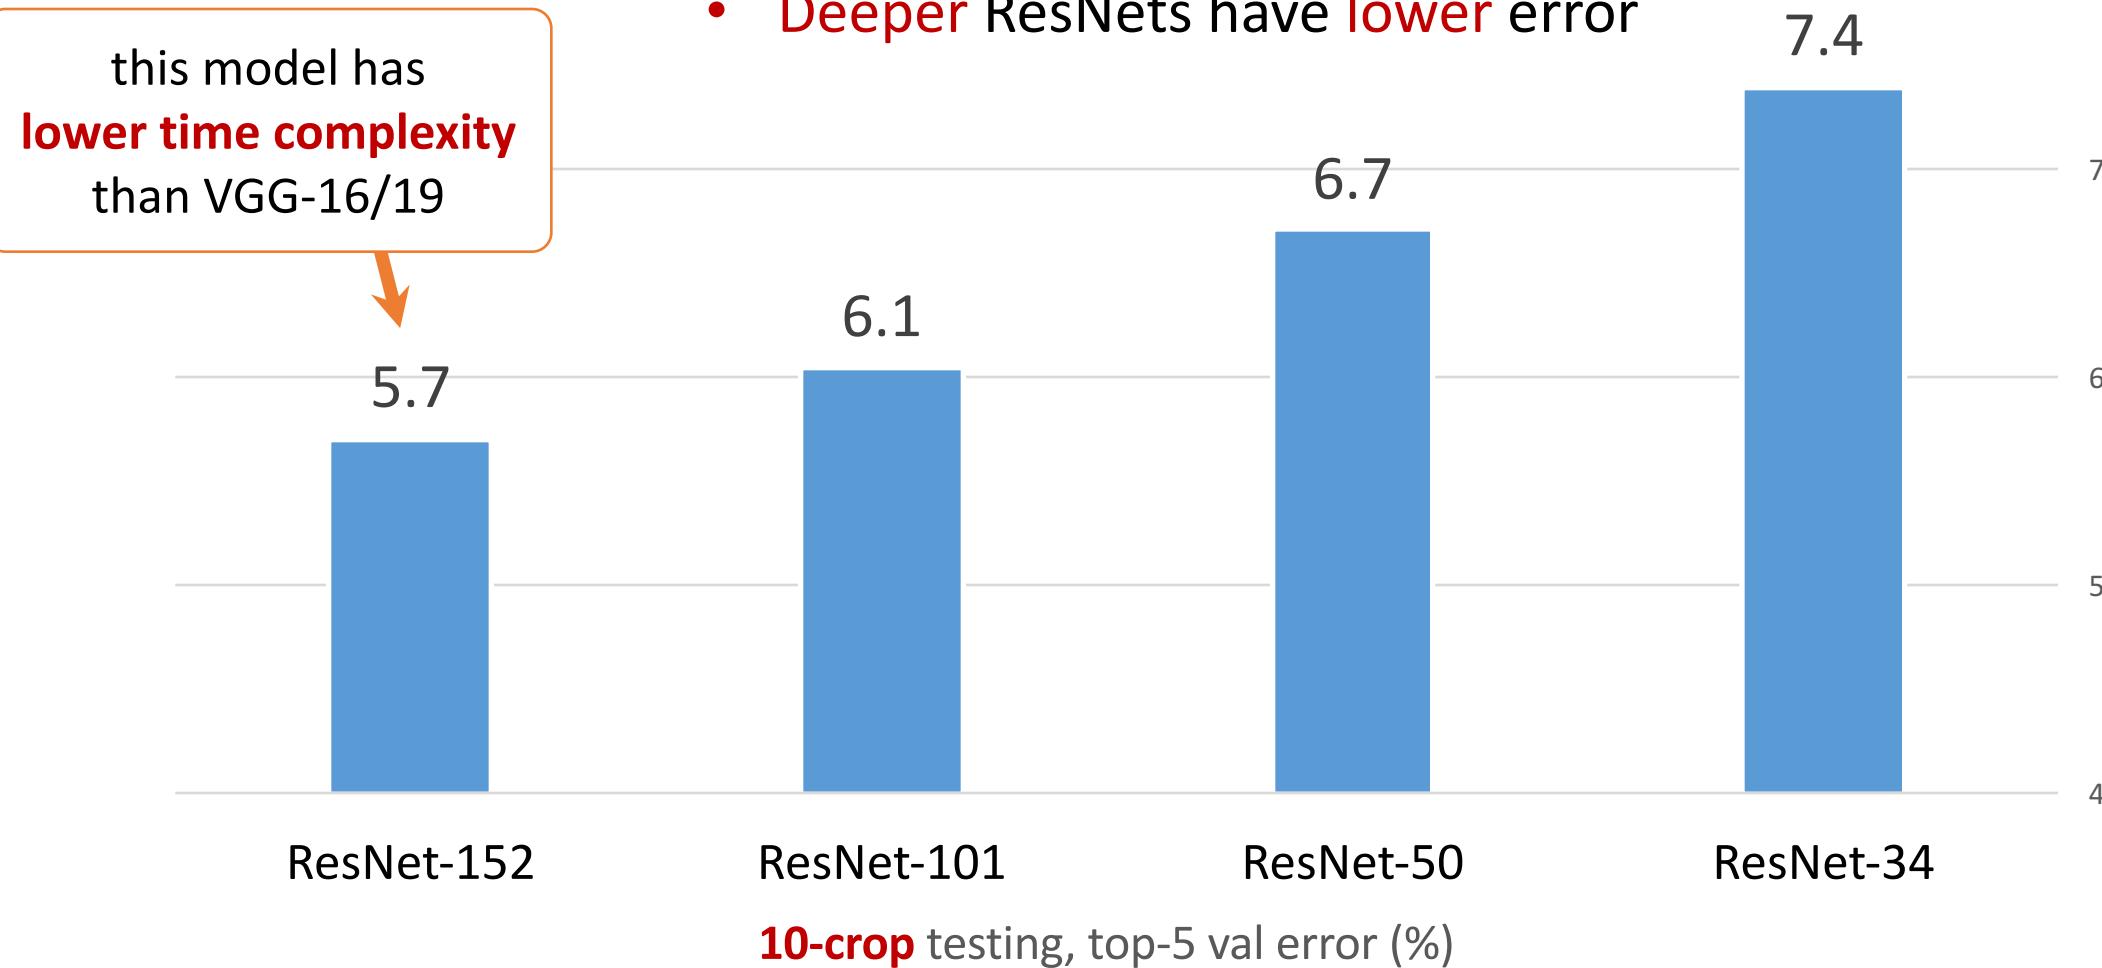


bottleneck  
(for ResNet-50/101/152)

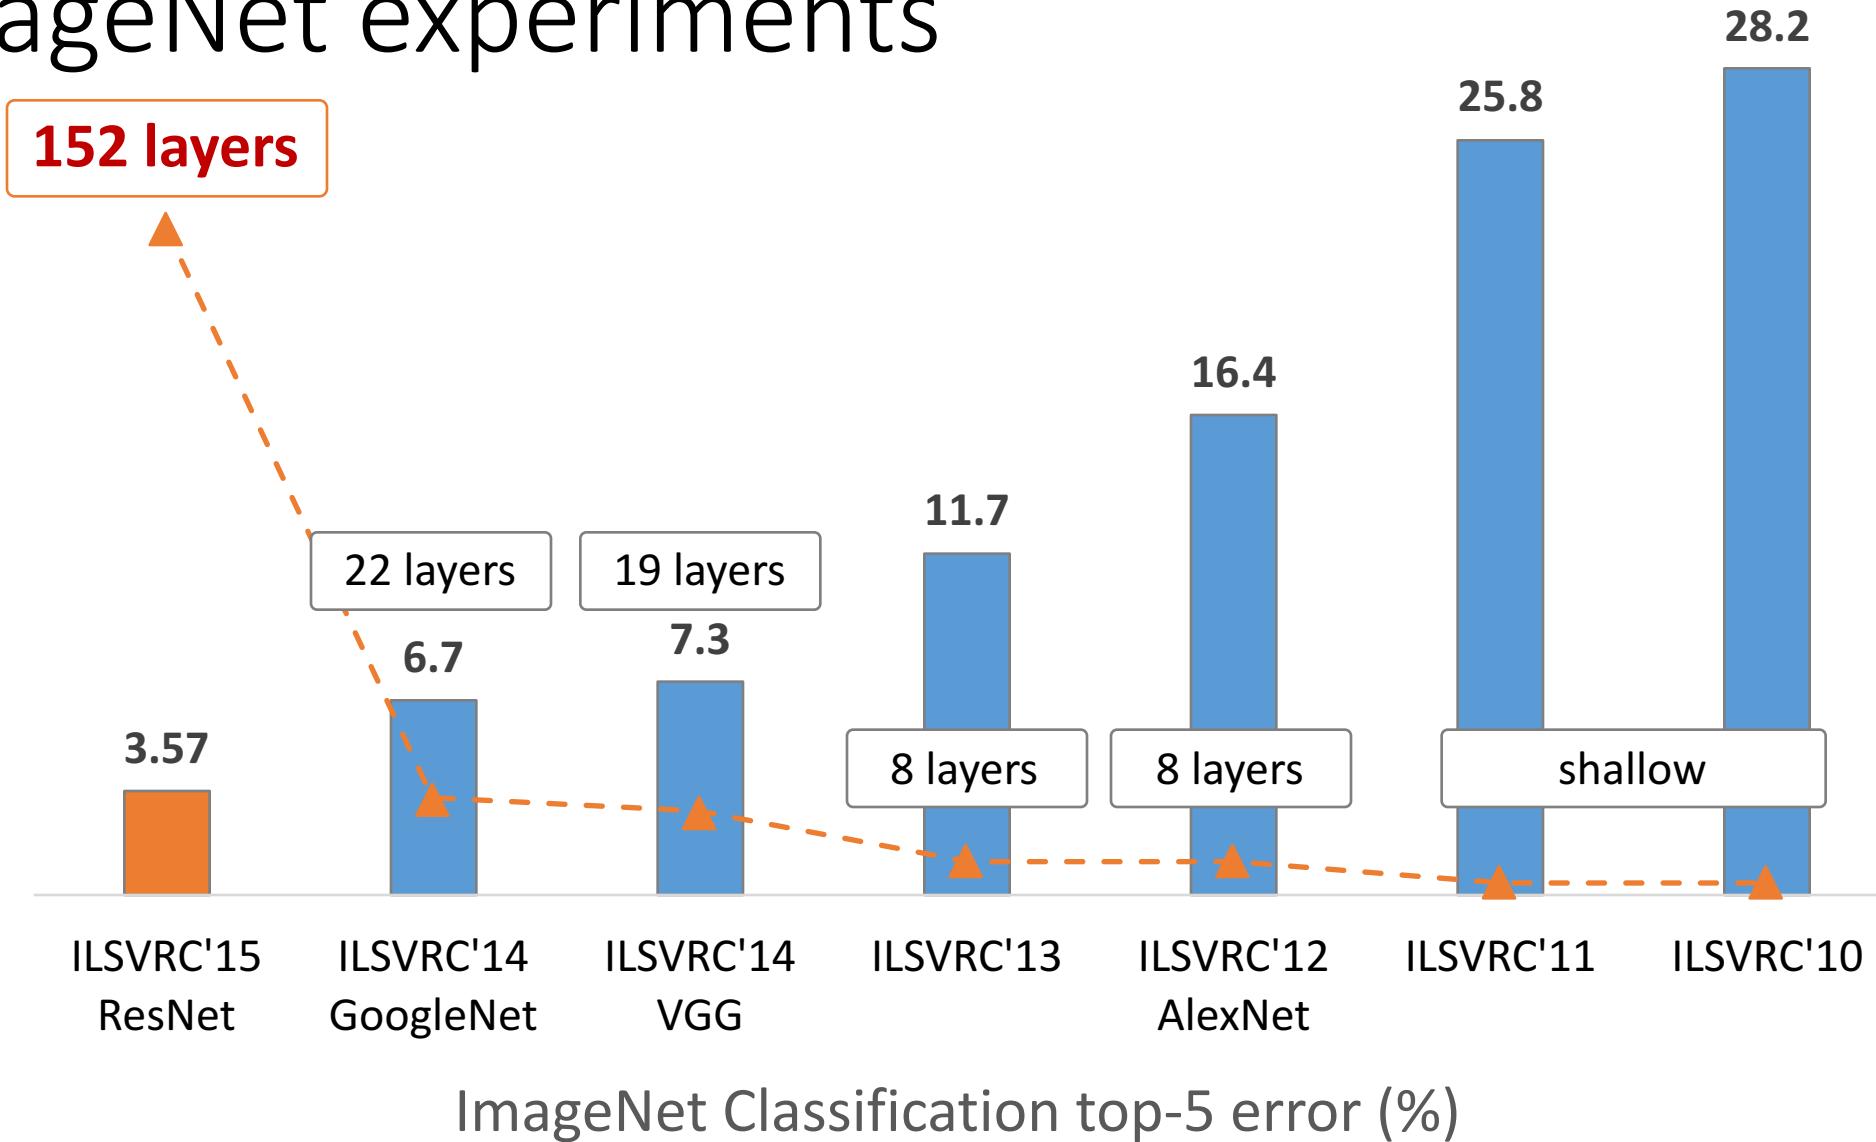
similar complexity

# ImageNet experiments

- Deeper ResNets have lower error



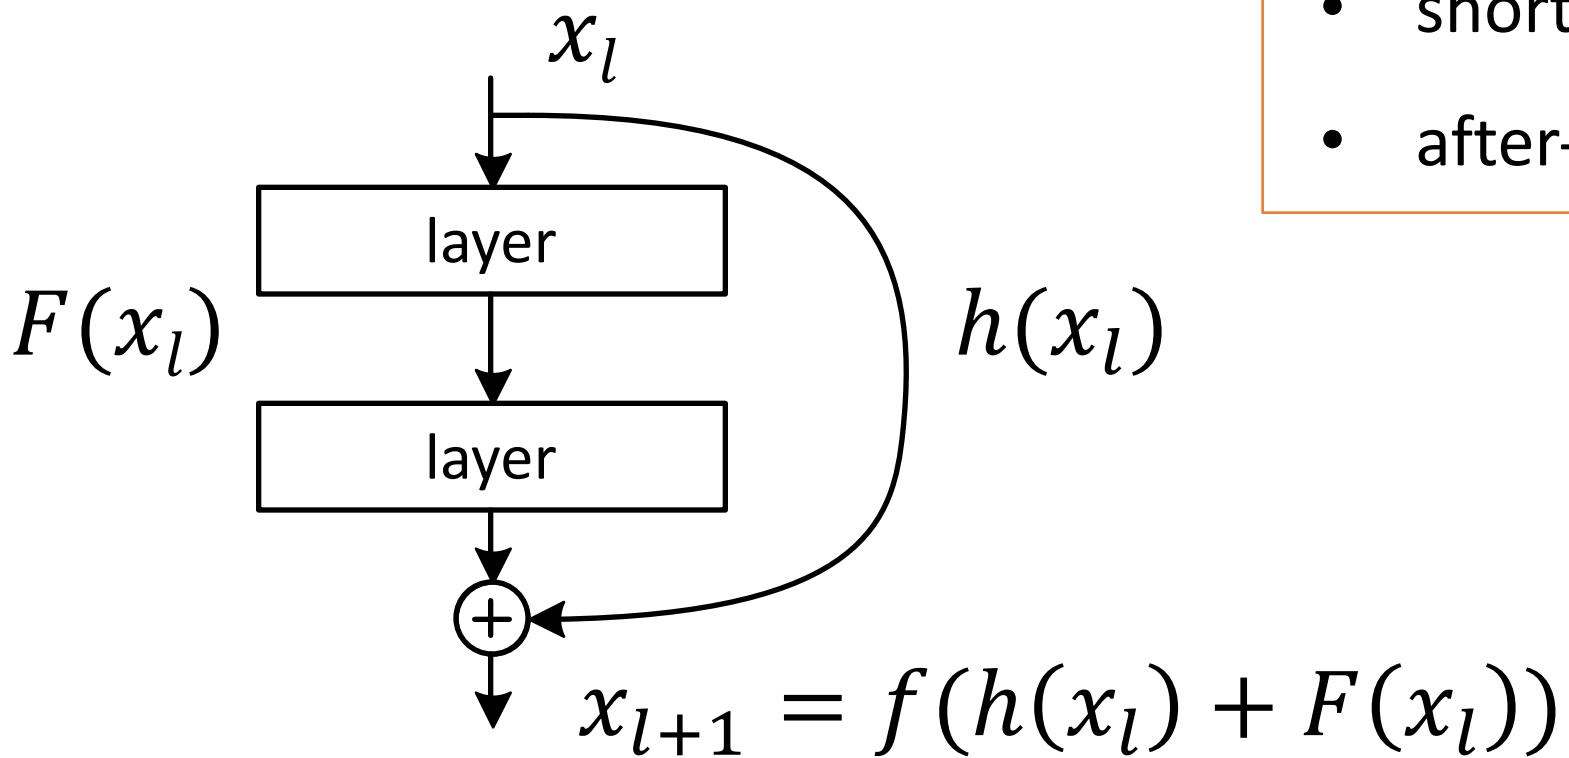
# ImageNet experiments



# On the Importance of Identity Mapping

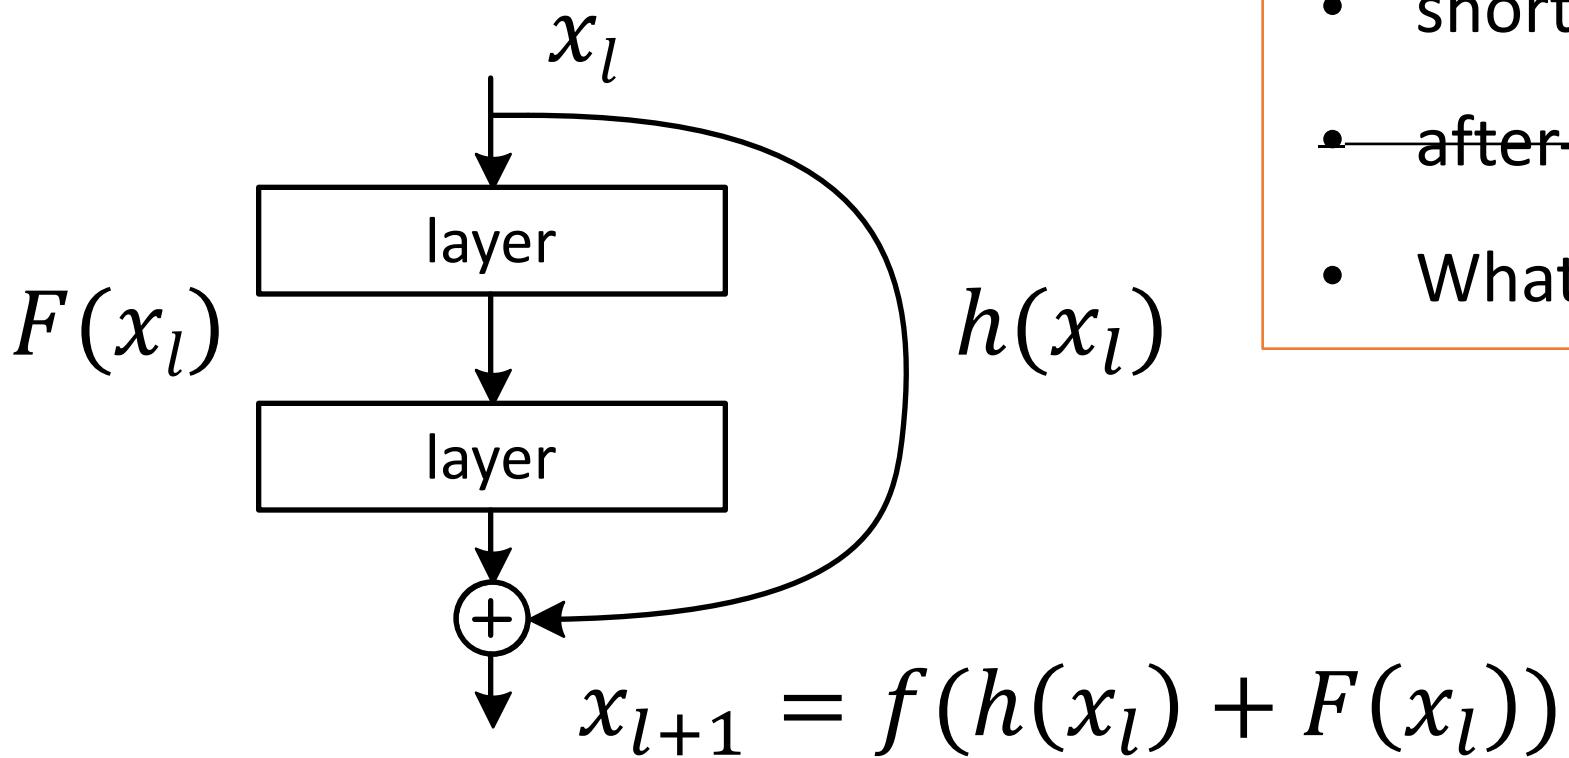
A Deeper Look at ResNets

# On identity mappings for optimization



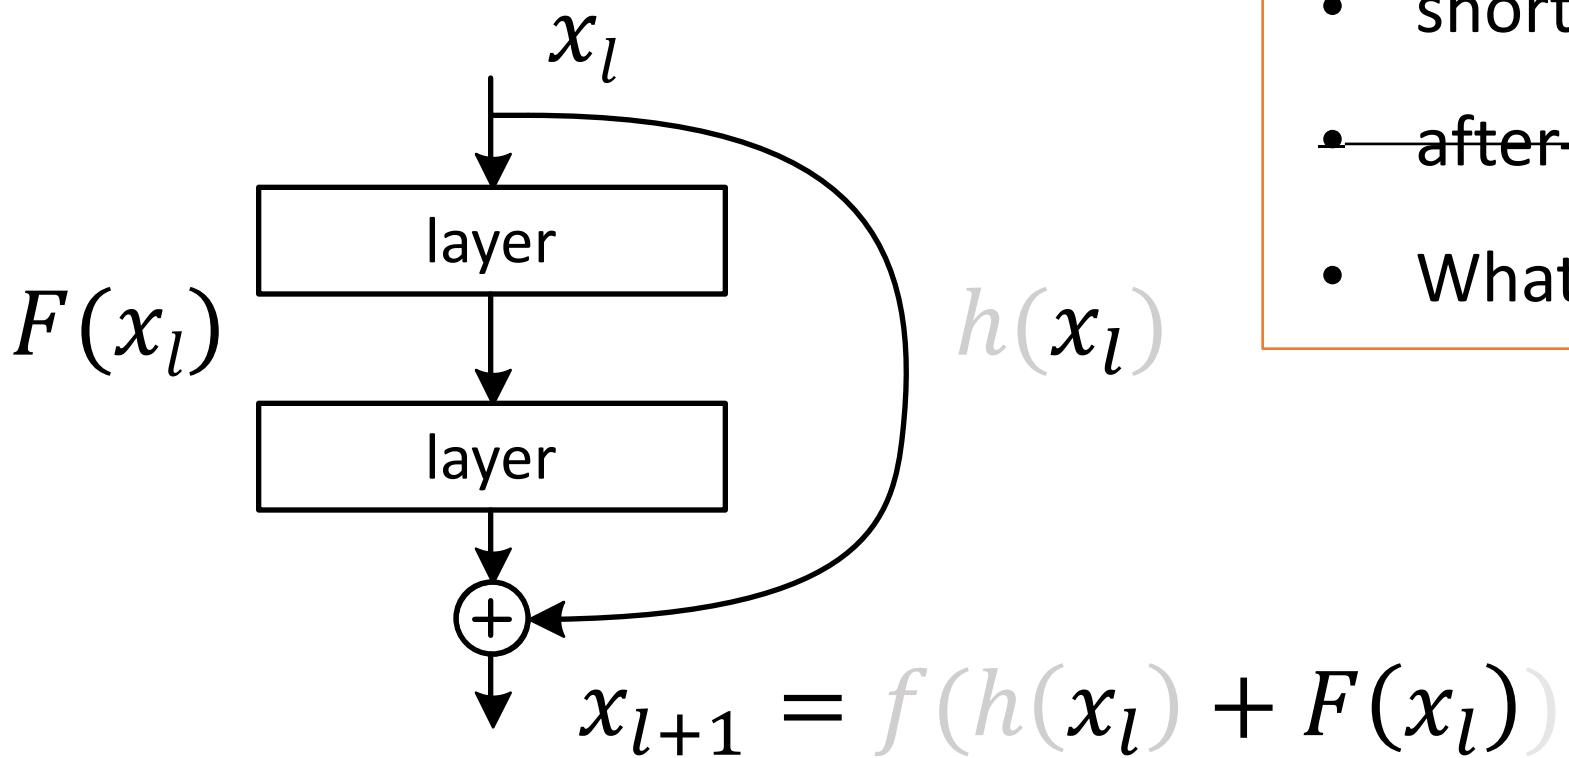
- shortcut mapping:  $h = \text{identity}$
- after-add mapping:  $f = \text{ReLU}$

# On identity mappings for optimization



- shortcut mapping:  $h = \text{identity}$
- after-add mapping:  $f = \text{ReLU}$
- What if  $f = \text{identity}$ ?

# On identity mappings for optimization



- shortcut mapping:  $h = \text{identity}$
- after-add mapping:  $f = \text{ReLU}$
- What if  $f = \text{identity}$ ?

# Very smooth forward propagation

$$x_{l+1} = x_l + F(x_l)$$



$$x_{l+2} = x_{l+1} + F(x_{l+1})$$

# Very smooth forward propagation

$$x_{l+1} = x_l + F(x_l)$$



$$x_{l+2} = x_{l+1} + F(x_{l+1})$$

$$x_{l+2} = x_l + F(x_l) + F(x_{l+1})$$

# Very smooth forward propagation

$$x_{l+1} = x_l + F(x_l)$$



$$x_{l+2} = x_{l+1} + F(x_{l+1})$$

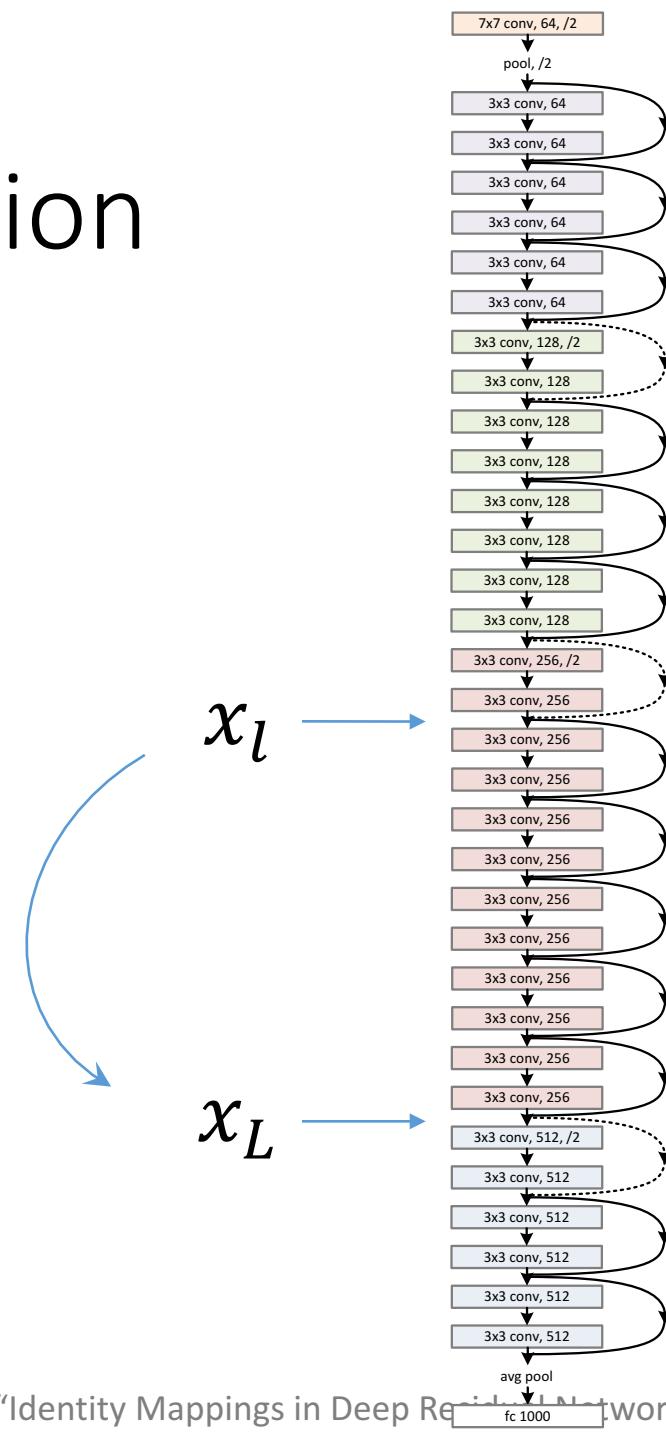
$$x_{l+2} = x_l + F(x_l) + F(x_{l+1})$$

$$x_L = x_l + \sum_{i=l}^{L-1} F(x_i)$$

# Very smooth forward propagation

$$x_L = x_l + \sum_{i=l}^{L-1} F(x_i)$$

- Any  $x_l$  is **directly** forward-prop to any  $x_L$ , plus **residual**.
- Any  $x_L$  is an **additive** outcome.
  - in contrast to **multiplicative**:  $x_L = \prod_{i=l}^{L-1} W_i x_l$



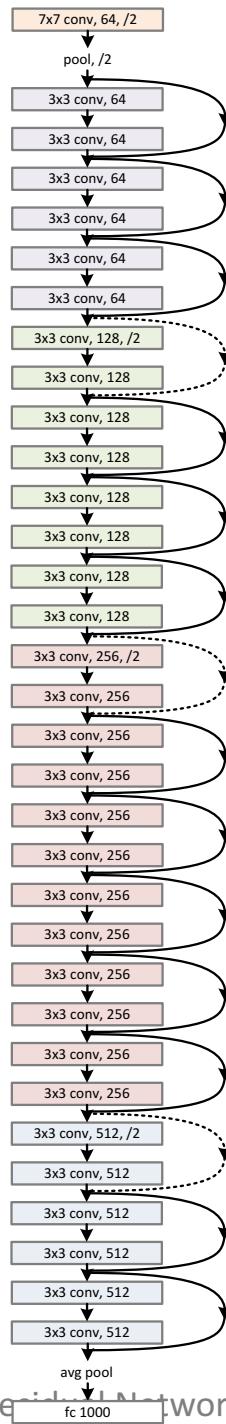
# Very smooth backward propagation

$$x_L = x_l + \sum_{i=l}^{L-1} F(x_i)$$



$$\frac{\partial E}{\partial x_l} = \frac{\partial E}{\partial x_L} \frac{\partial x_L}{\partial x_l} = \frac{\partial E}{\partial x_L} \left( 1 + \frac{\partial}{\partial x_l} \sum_{i=1}^{L-1} F(x_i) \right)$$

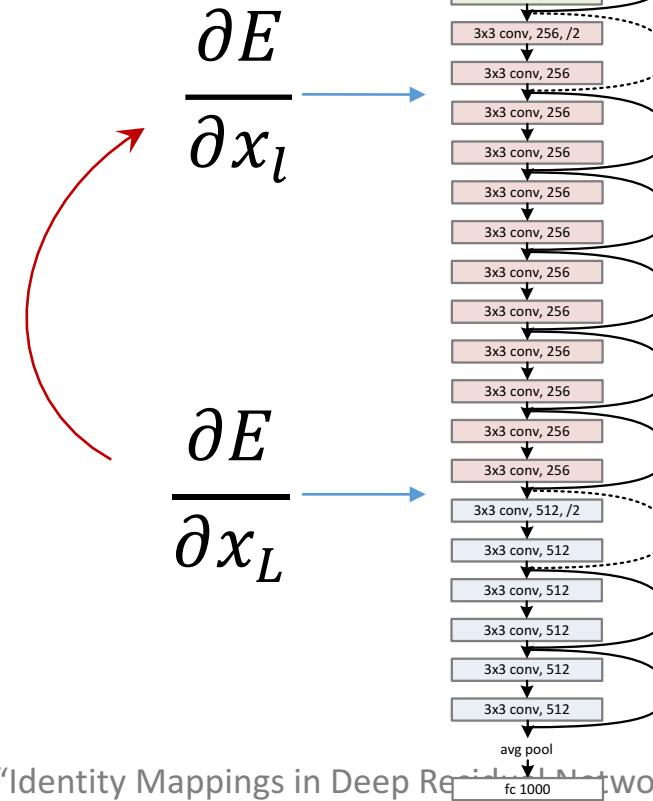
$$\frac{\partial E}{\partial x_l} \quad \frac{\partial E}{\partial x_L}$$



# Very smooth backward propagation

$$\frac{\partial E}{\partial x_l} = \frac{\partial E}{\partial x_L} \left( 1 + \frac{\partial}{\partial x_l} \sum_{i=1}^{L-1} F(x_i) \right)$$

- Any  $\frac{\partial E}{\partial x_L}$  is **directly** back-prop to any  $\frac{\partial E}{\partial x_l}$ , plus **residual**.
- Any  $\frac{\partial E}{\partial x_l}$  is **additive**; unlikely to vanish
  - in contrast to **multiplicative**:  $\frac{\partial E}{\partial x_l} = \prod_{i=l}^{L-1} W_i \frac{\partial E}{\partial x_L}$



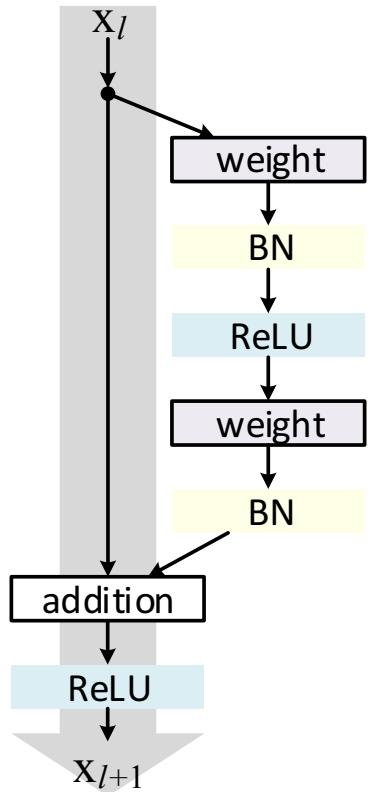
# Residual for every layer

forward:  $x_L = x_l + \sum_{i=l}^{L-1} F(x_i)$

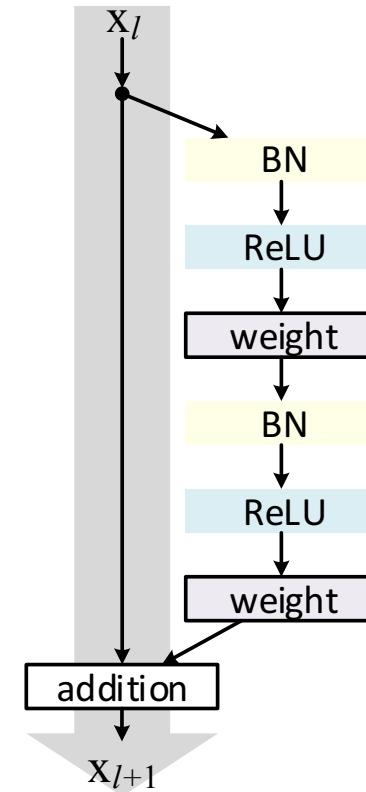
Enabled by:

- shortcut mapping:  $h = \text{identity}$
- after-add mapping:  $f = \text{identity}$

backward:  $\frac{\partial E}{\partial x_l} = \frac{\partial E}{\partial x_L} \left( 1 + \frac{\partial}{\partial x_l} \sum_{i=1}^{L-1} F(x_i) \right)$

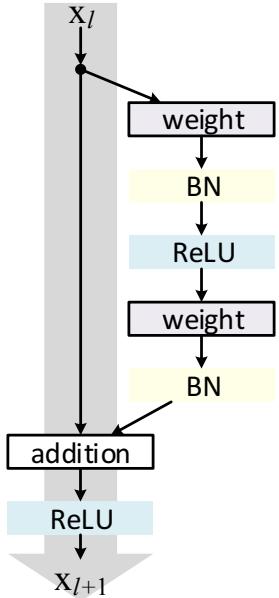
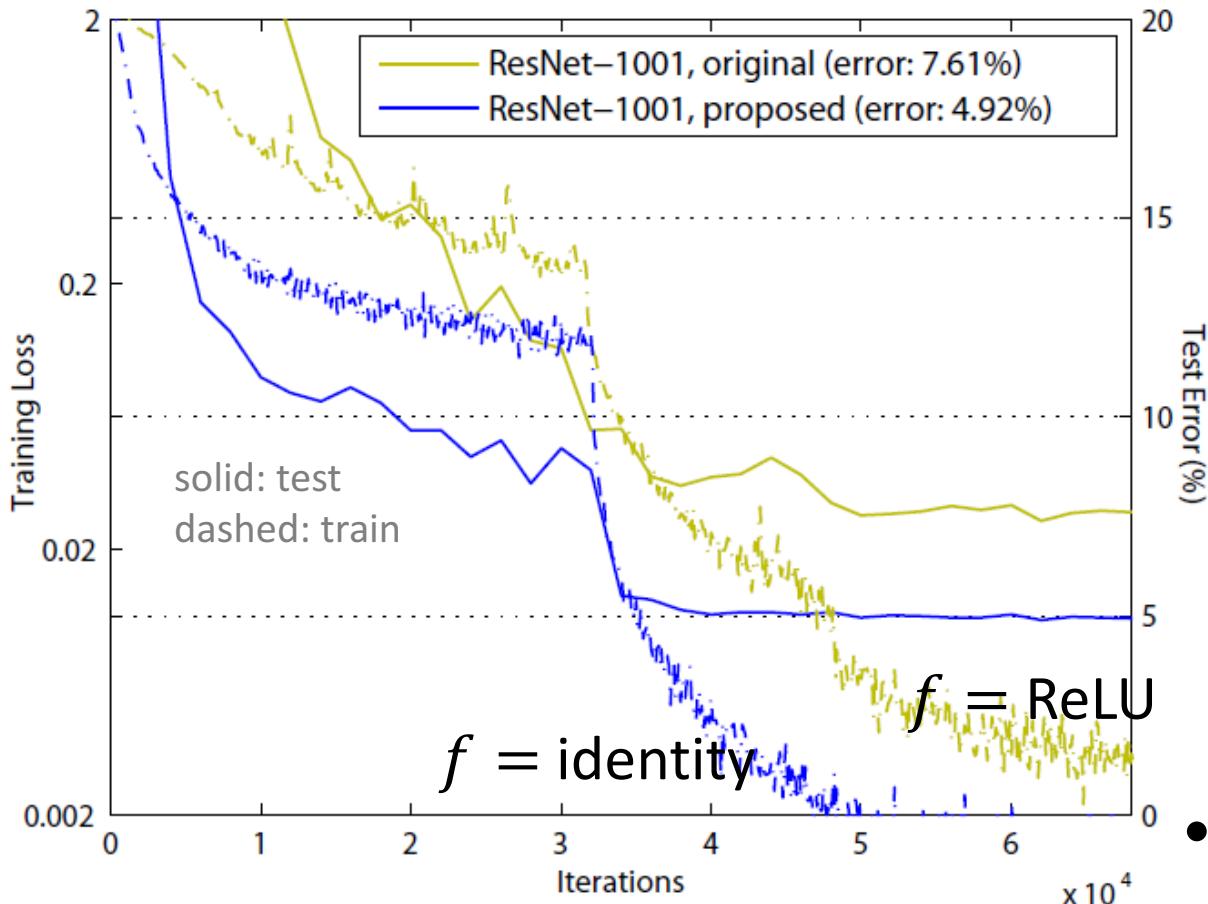


original ResNet



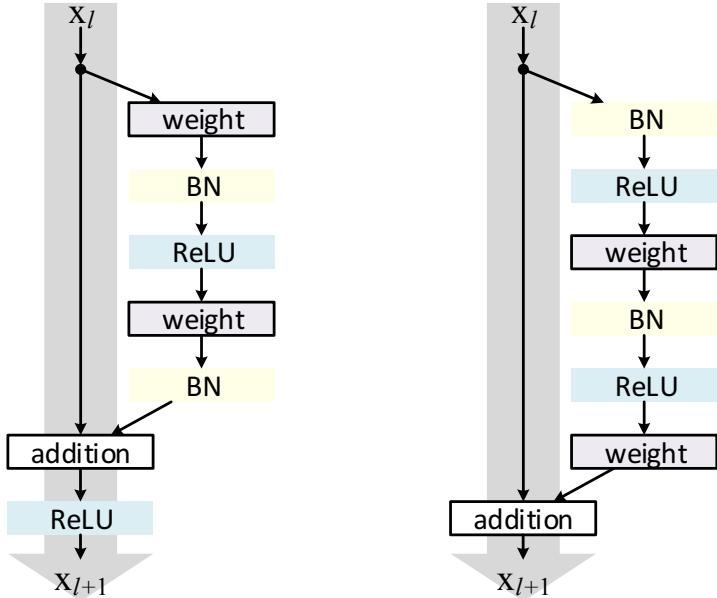
**pre-activation**  
ResNet

# 1001-layer ResNets on CIFAR-10



$$f = \text{ReLU}$$

$$f = \text{identity}$$



- ReLU could also block prop when there are 1000 layers
- pre-activation design eases optimization (and improves generalization; see paper)

# Comparisons on CIFAR-10/100

CIFAR-10

method	error (%)
NIN	8.81
DSN	8.22
FitNet	8.39
Highway	7.72
ResNet-110 (1.7M)	6.61
ResNet-1202 (19.4M)	7.93
ResNet-164, pre-activation (1.7M)	5.46
<b>ResNet-1001</b> , pre-activation (10.2M)	<b>4.92</b> ( $4.89 \pm 0.14$ )

CIFAR-100

method	error (%)
NIN	35.68
DSN	34.57
FitNet	35.04
Highway	32.39
ResNet-164 (1.7M)	25.16
ResNet-1001 (10.2M)	27.82
ResNet-164, pre-activation (1.7M)	24.33
<b>ResNet-1001</b> , pre-activation (10.2M)	<b>22.71</b> ( $22.68 \pm 0.22$ )

\*all based on moderate augmentation

# ImageNet Experiments

ImageNet single-crop (320x320) val error

method	data augmentation	top-1 error (%)	top-5 error (%)
ResNet-152, original	scale	21.3	5.5
ResNet-152, pre-activation	scale	21.1	5.5
ResNet-200, original	scale	21.8	6.0
ResNet-200, pre-activation	scale	<b>20.7</b>	<b>5.3</b>
ResNet-200, pre-activation	scale + aspect ratio	<b>20.1*</b>	<b>4.8*</b>

\* <https://github.com/facebook/fb.resnet.torch/tree/master/pretrained#notes>  
**training code and models available.**

# From Classification to Detection

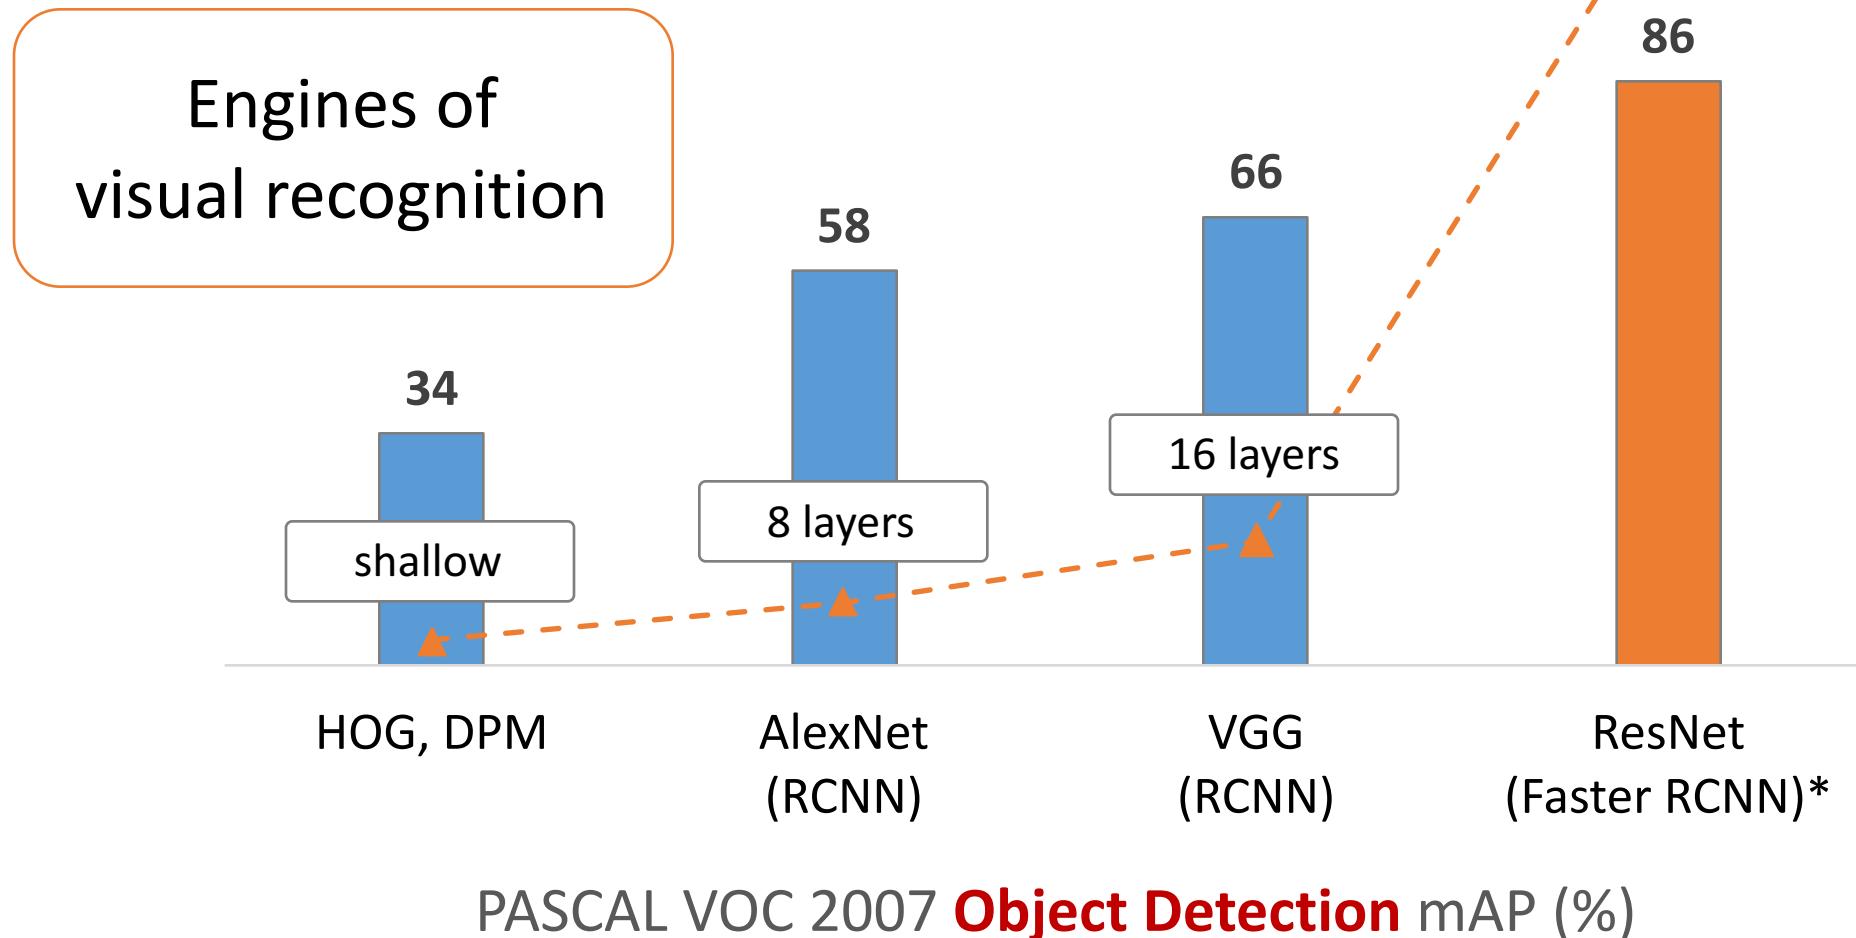
*“Features matter”*

*“Features matter.”* (quote [Girshick et al. 2014], the R-CNN paper)

task	2nd-place winner	ResNets	margin (relative)
ImageNet Localization <small>(top-5 error)</small>	12.0	9.0	<b>27%</b>
ImageNet Detection <small>(mAP@.5)</small>	53.6	62.1	<b>16%</b>
COCO Detection <small>(mAP@.5:.95)</small>	33.5	37.3	<b>11%</b>
COCO Segmentation <small>(mAP@.5:.95)</small>	25.1	28.2	<b>12%</b>

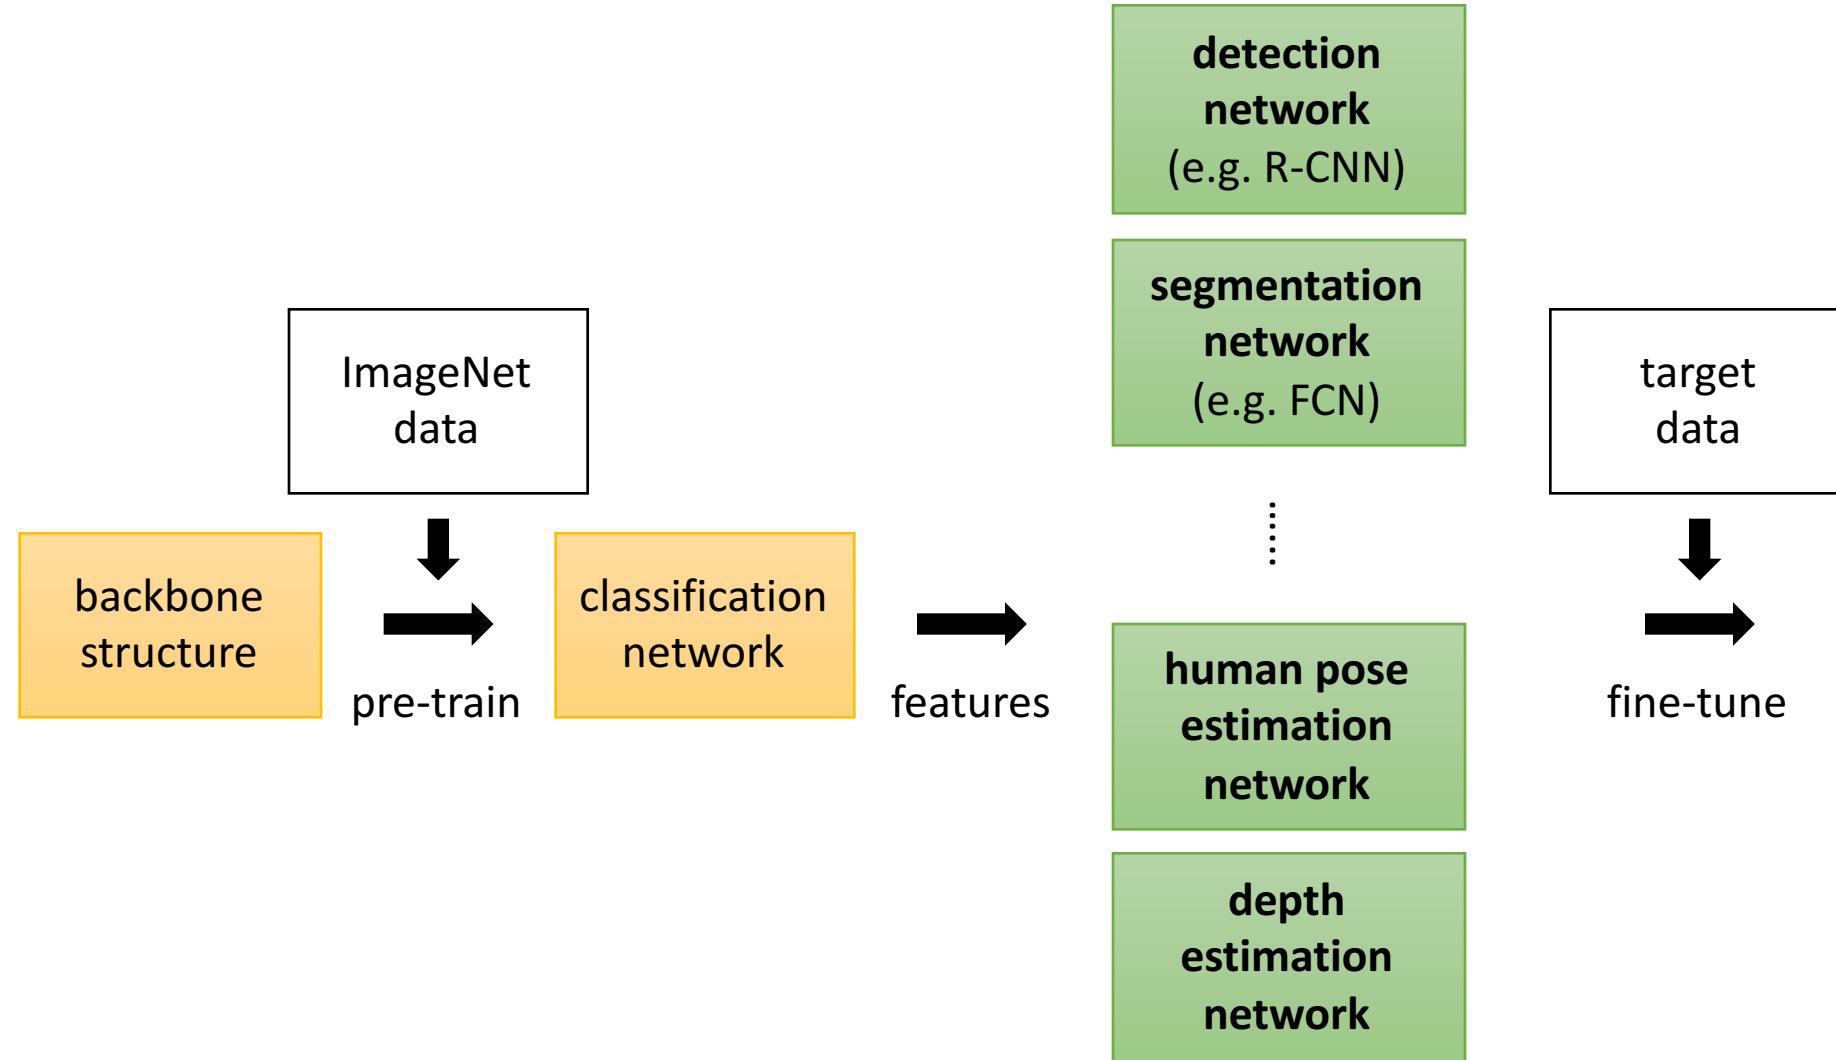
- Our results are all based on **ResNet-101**
- Deeper features are **well transferrable**

# Revolution of Depth



\*w/ other improvements & more data

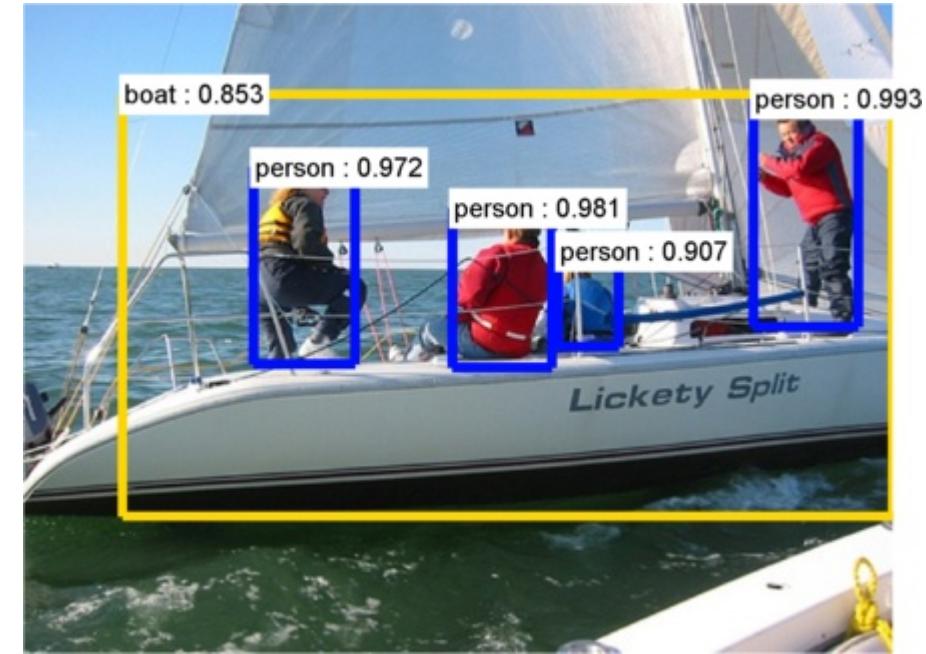
# Deep Learning for Computer Vision



# Example: Object Detection



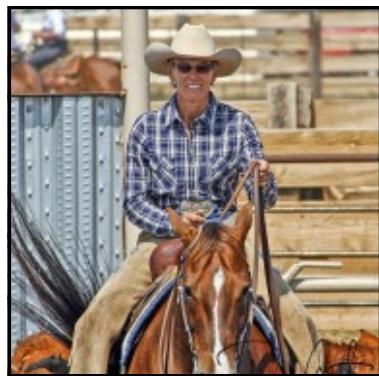
Image Classification  
(what?)



Object Detection  
(what + where?)

# Object Detection: R-CNN

figure credit: R. Girshick et al.

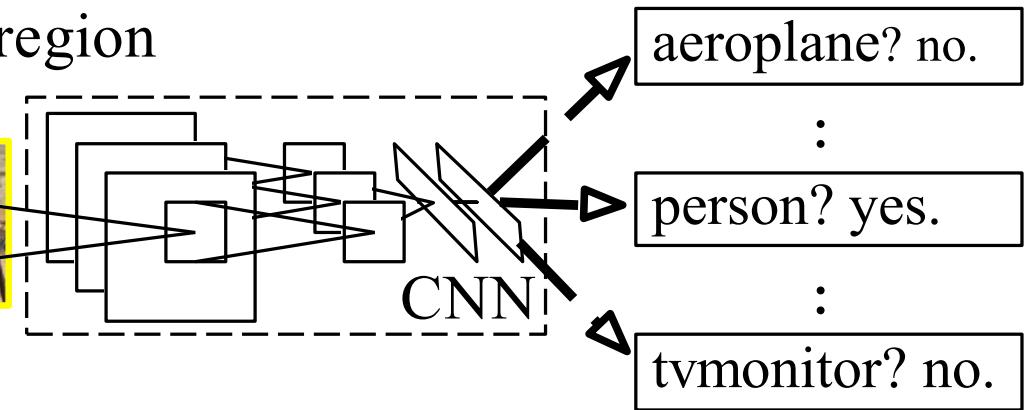


input image



region proposals  
~2,000

warped region



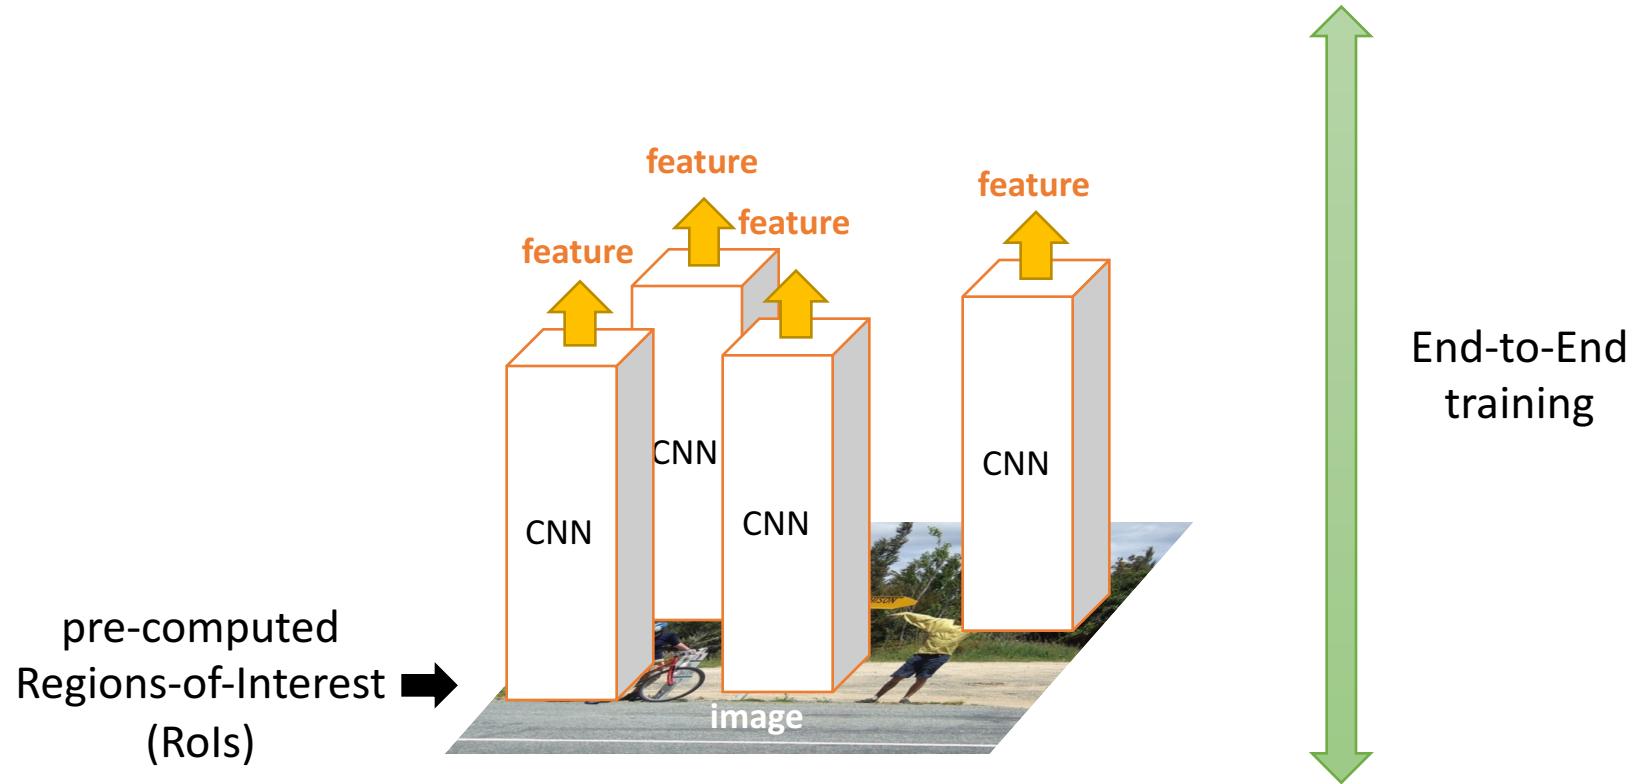
1 CNN for each region

classify regions

Region-based **CNN** pipeline

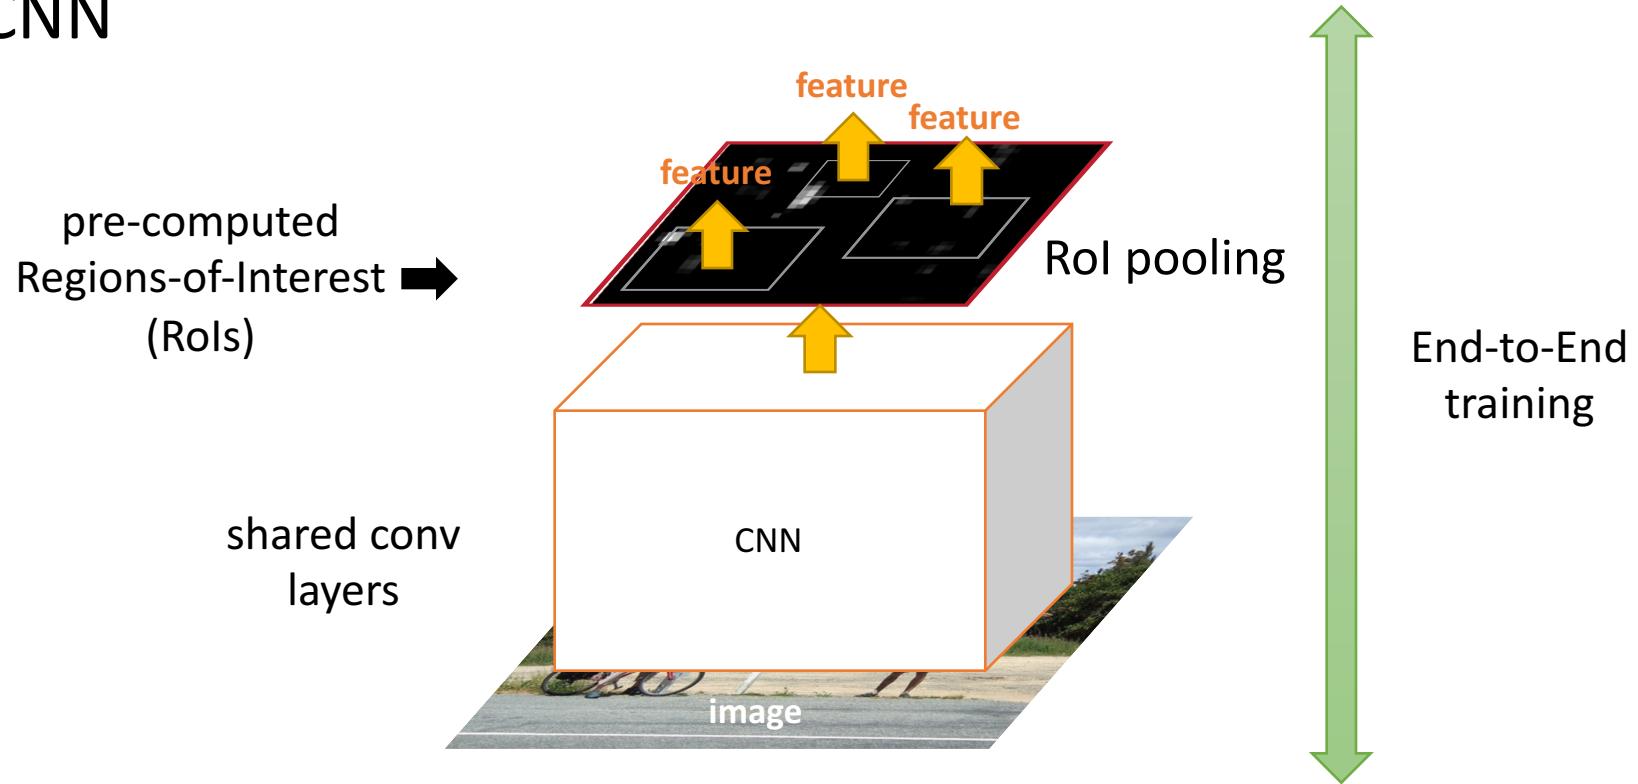
# Object Detection: R-CNN

- R-CNN



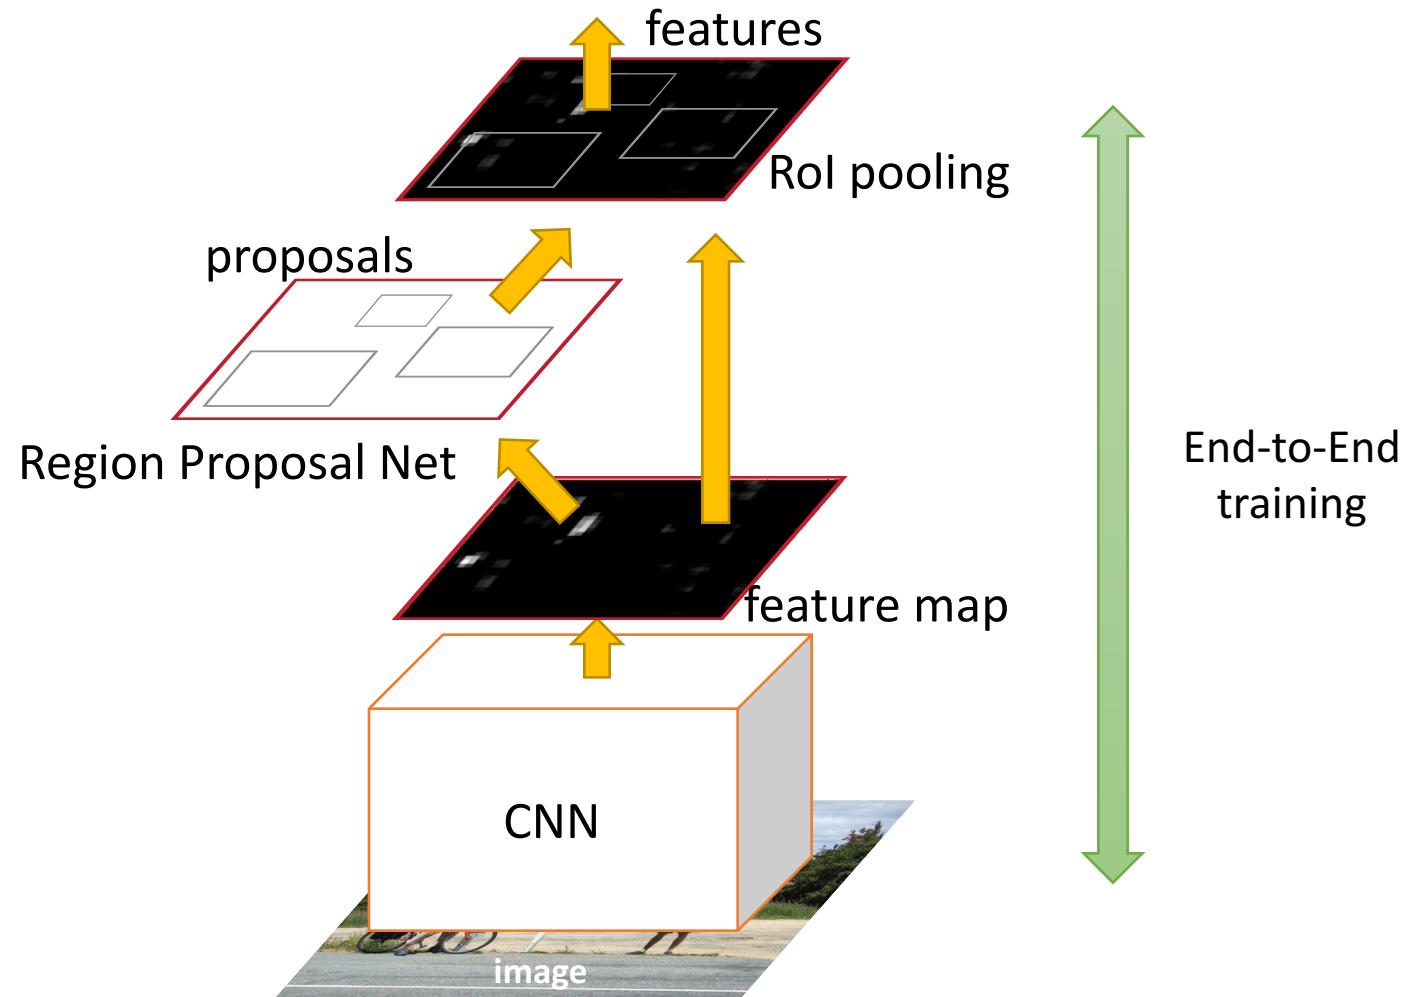
# Object Detection: Fast R-CNN

- Fast R-CNN

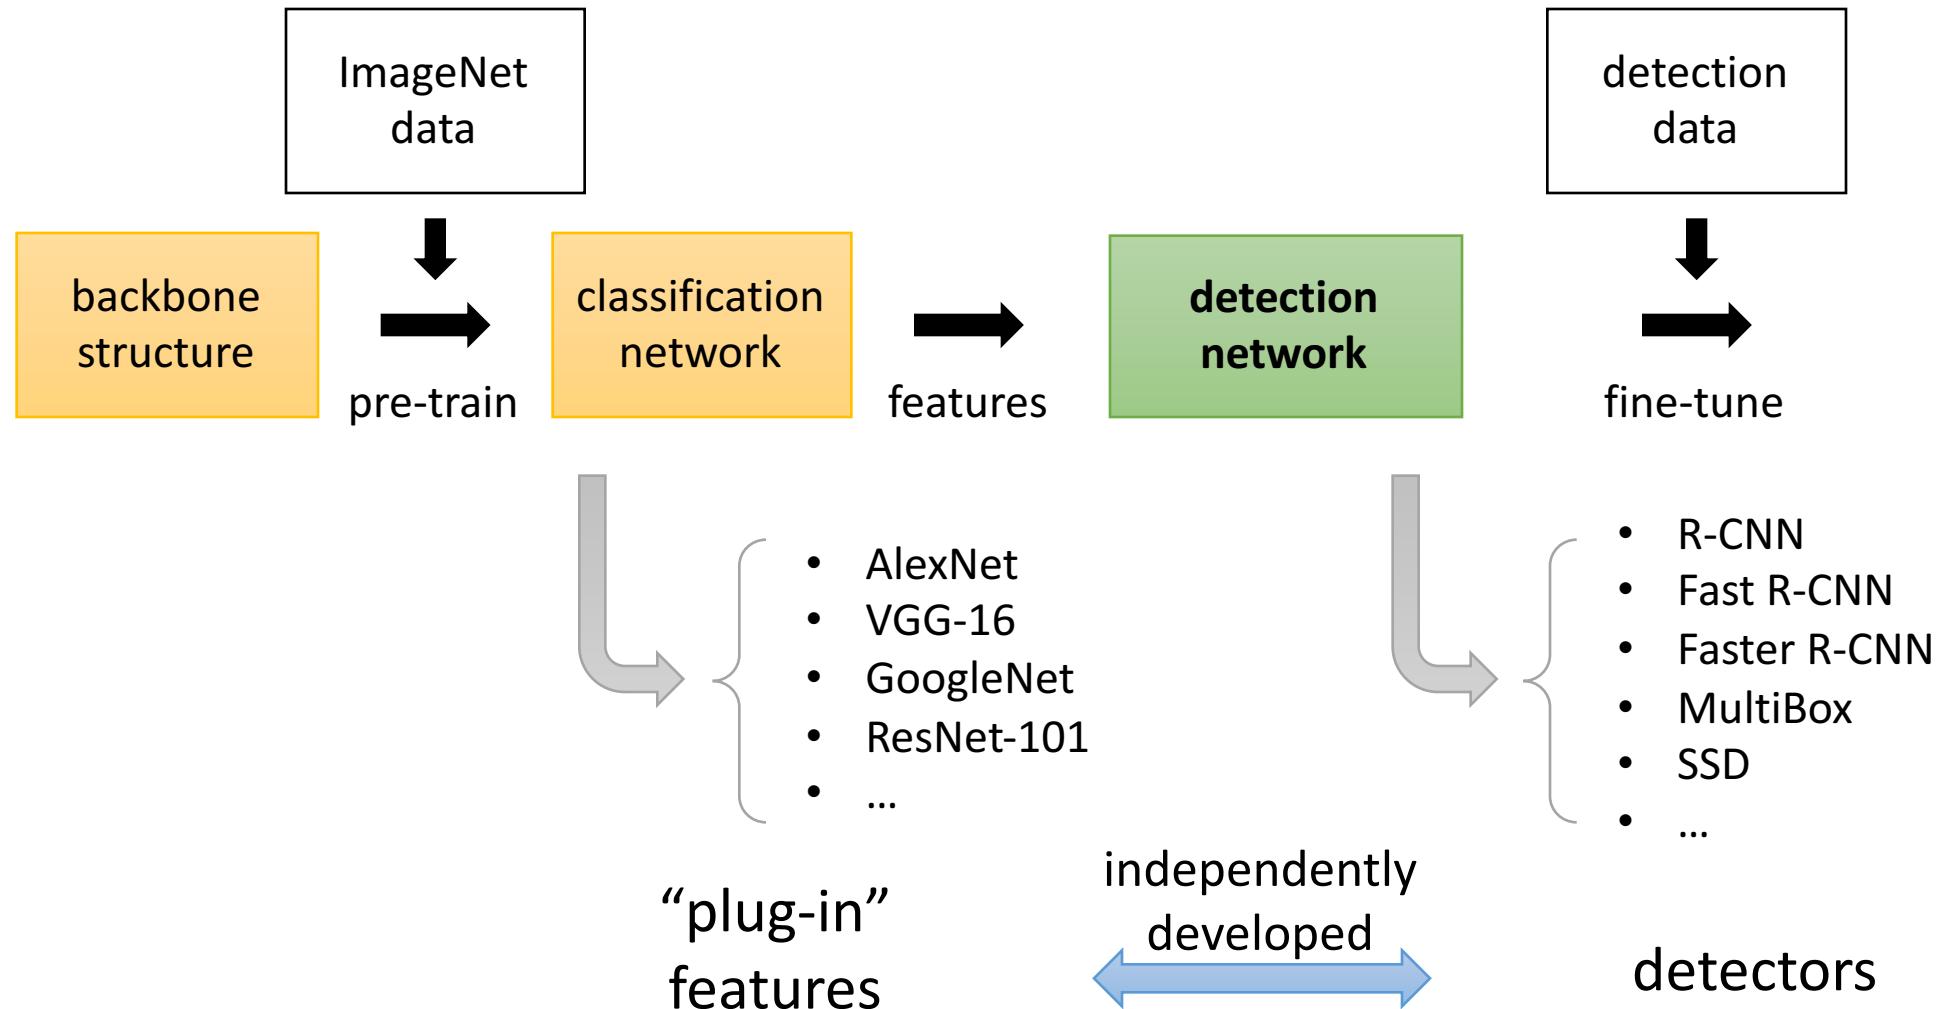


# Object Detection: Faster R-CNN

- Faster R-CNN
  - Solely based on CNN
  - No external modules
  - Each step is end-to-end



# Object Detection



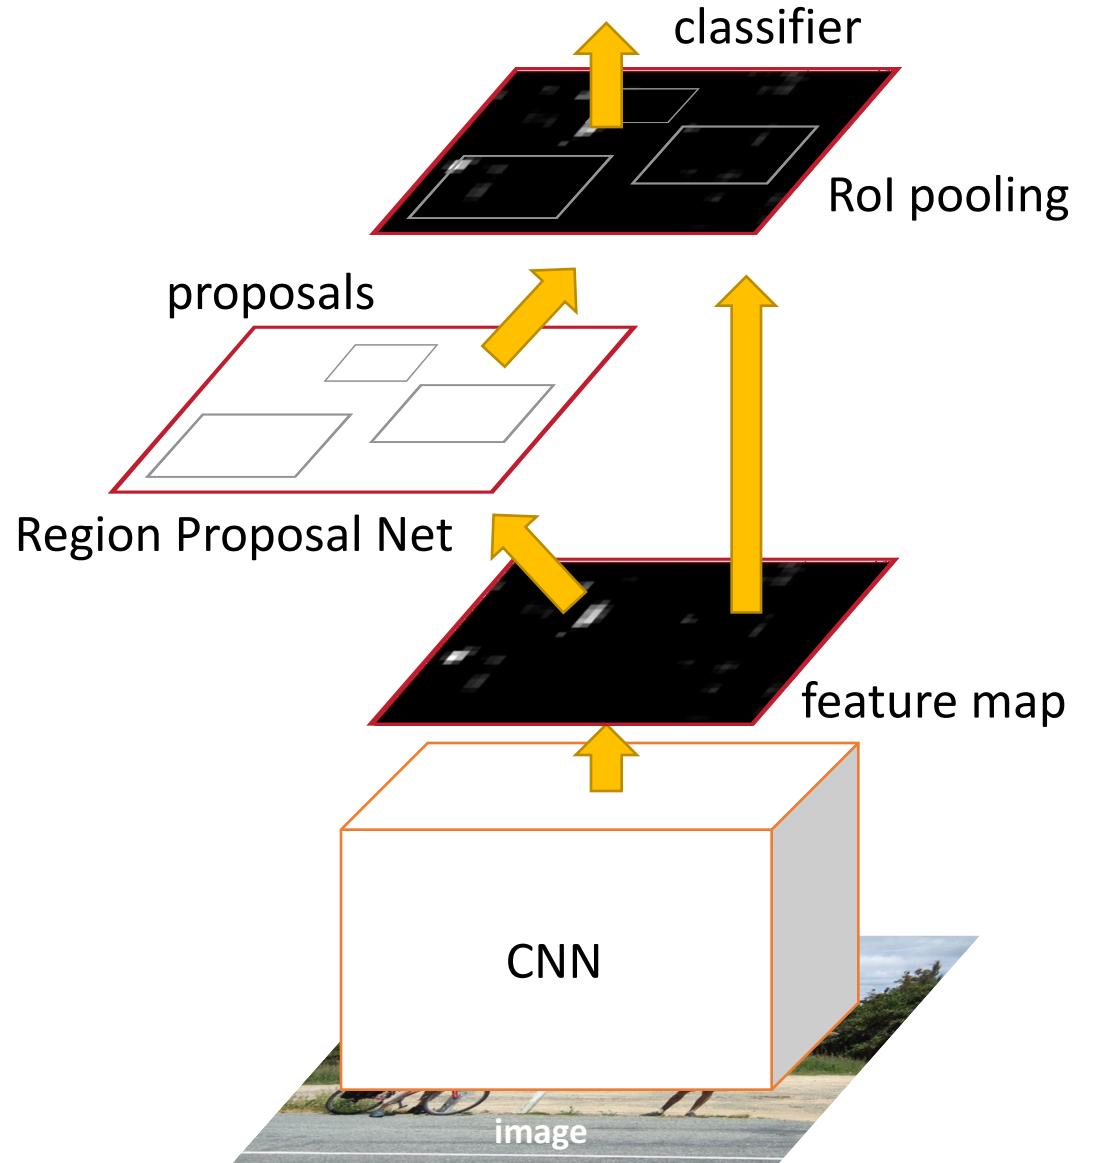
# Object Detection

- Simply “Faster R-CNN + ResNet”

Faster R-CNN baseline	mAP@.5	mAP@.5:.95
VGG-16	41.5	21.5
ResNet-101	<b>48.4</b>	<b>27.2</b>

coco detection results

**ResNet-101 has 28% relative gain  
vs VGG-16**

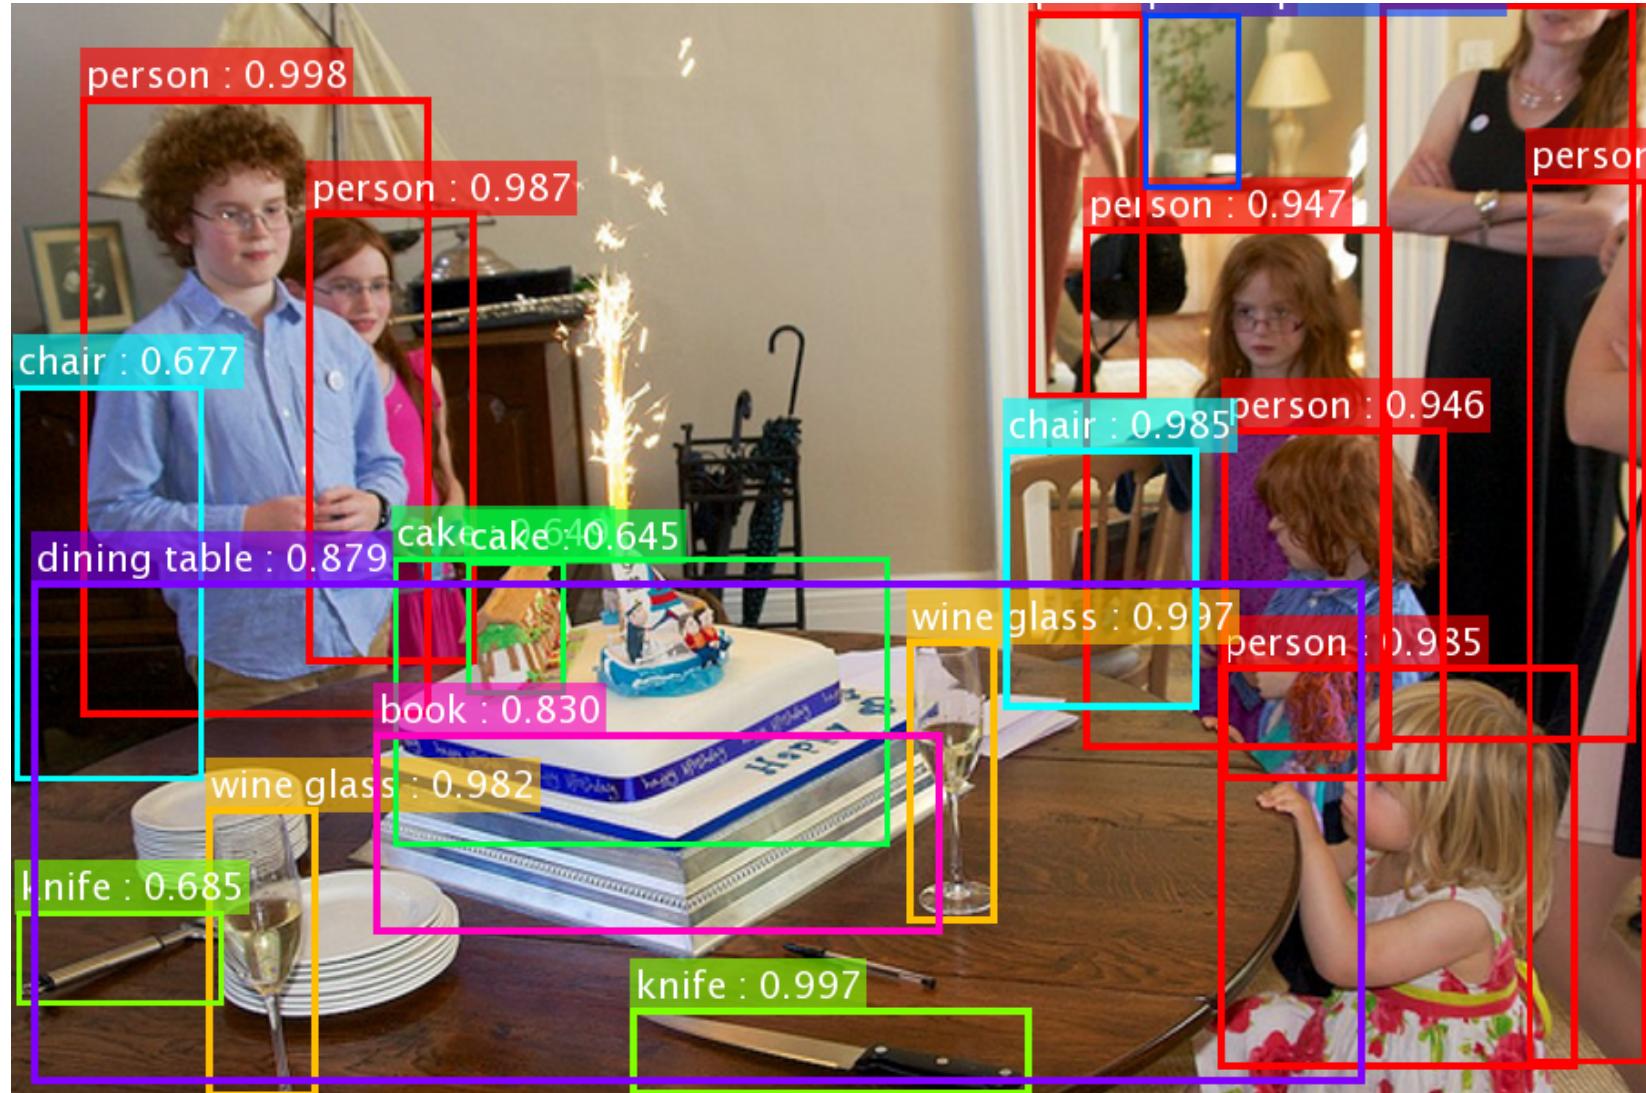


Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. “Deep Residual Learning for Image Recognition”. CVPR 2016.

Shaoqing Ren, Kaiming He, Ross Girshick, & Jian Sun. “Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks”. NIPS 2015.

# Object Detection

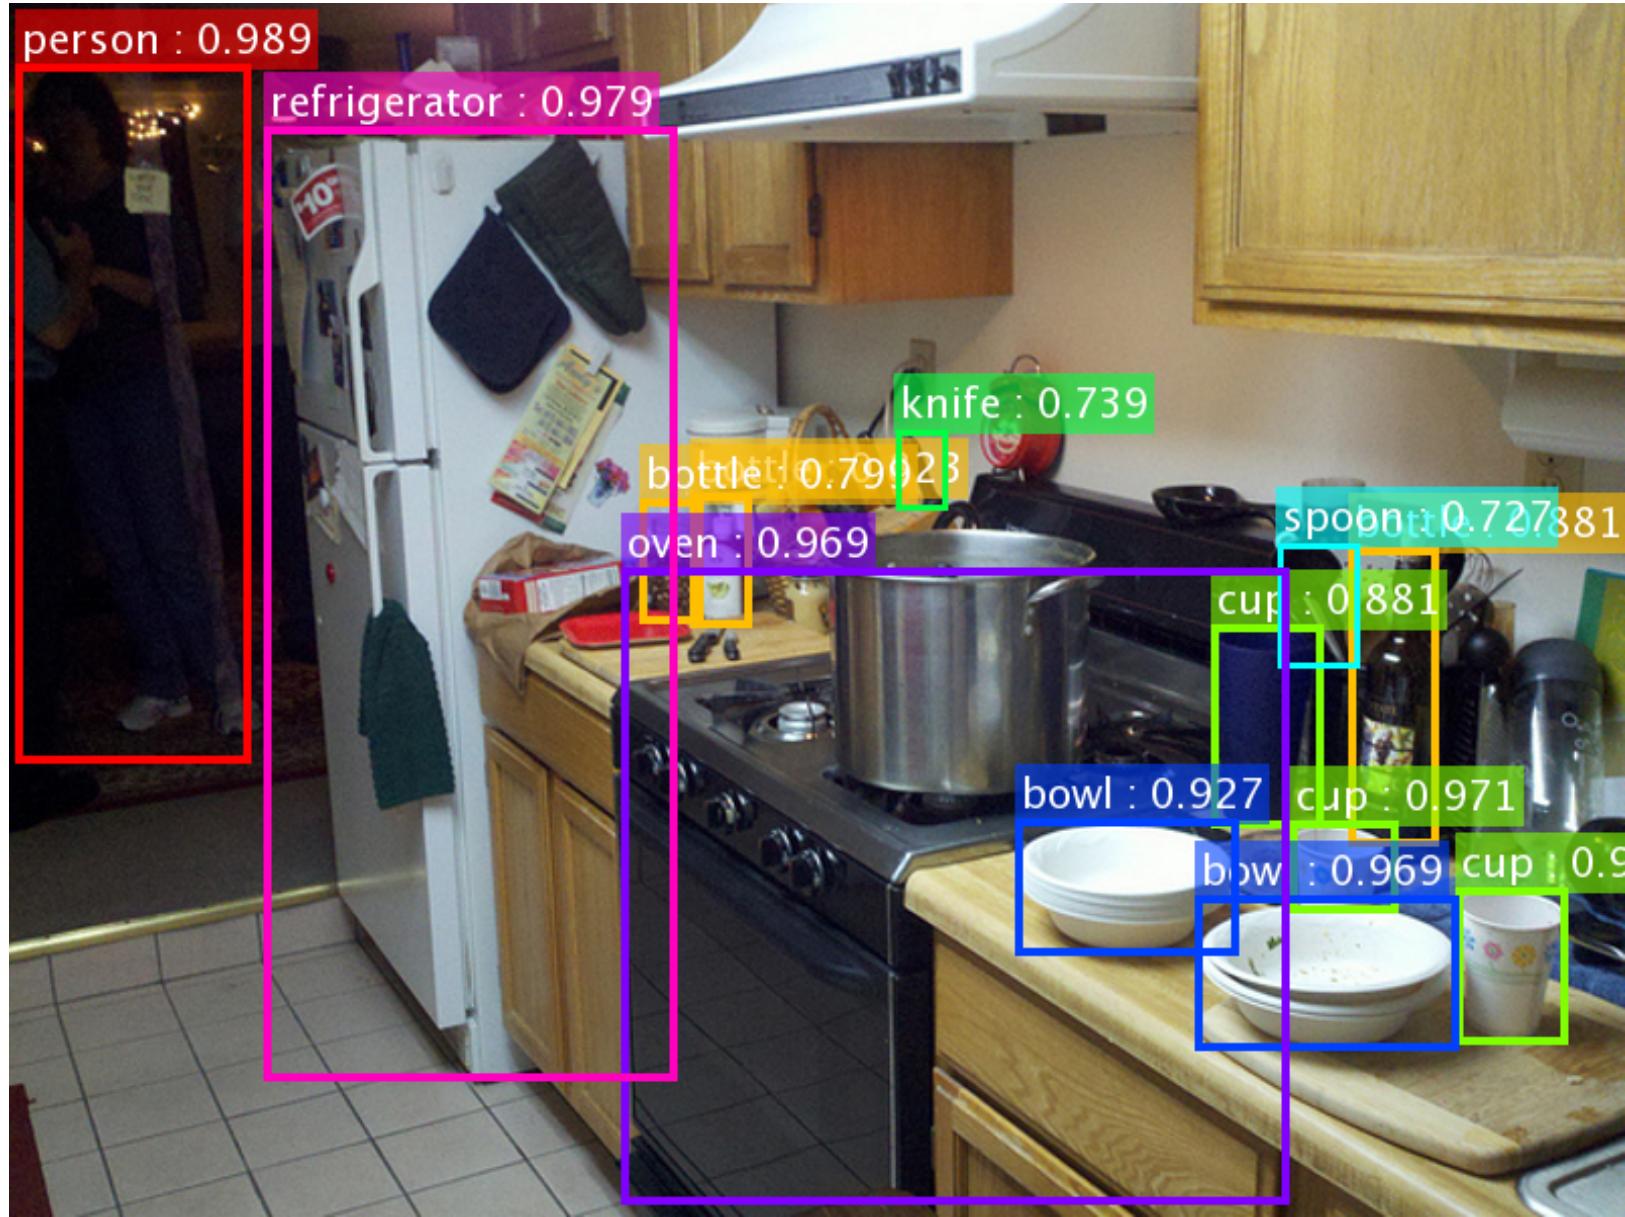
- RPN **learns** proposals by extremely deep nets
  - We use **only 300 proposals** (no hand-designed proposals)
- Add components:
  - Iterative localization
  - Context modeling
  - Multi-scale testing
- All components are based on CNN features; all steps are end-to-end
- All benefit **more** from **deeper** features – cumulative gains!



## ResNet's object detection result on COCO

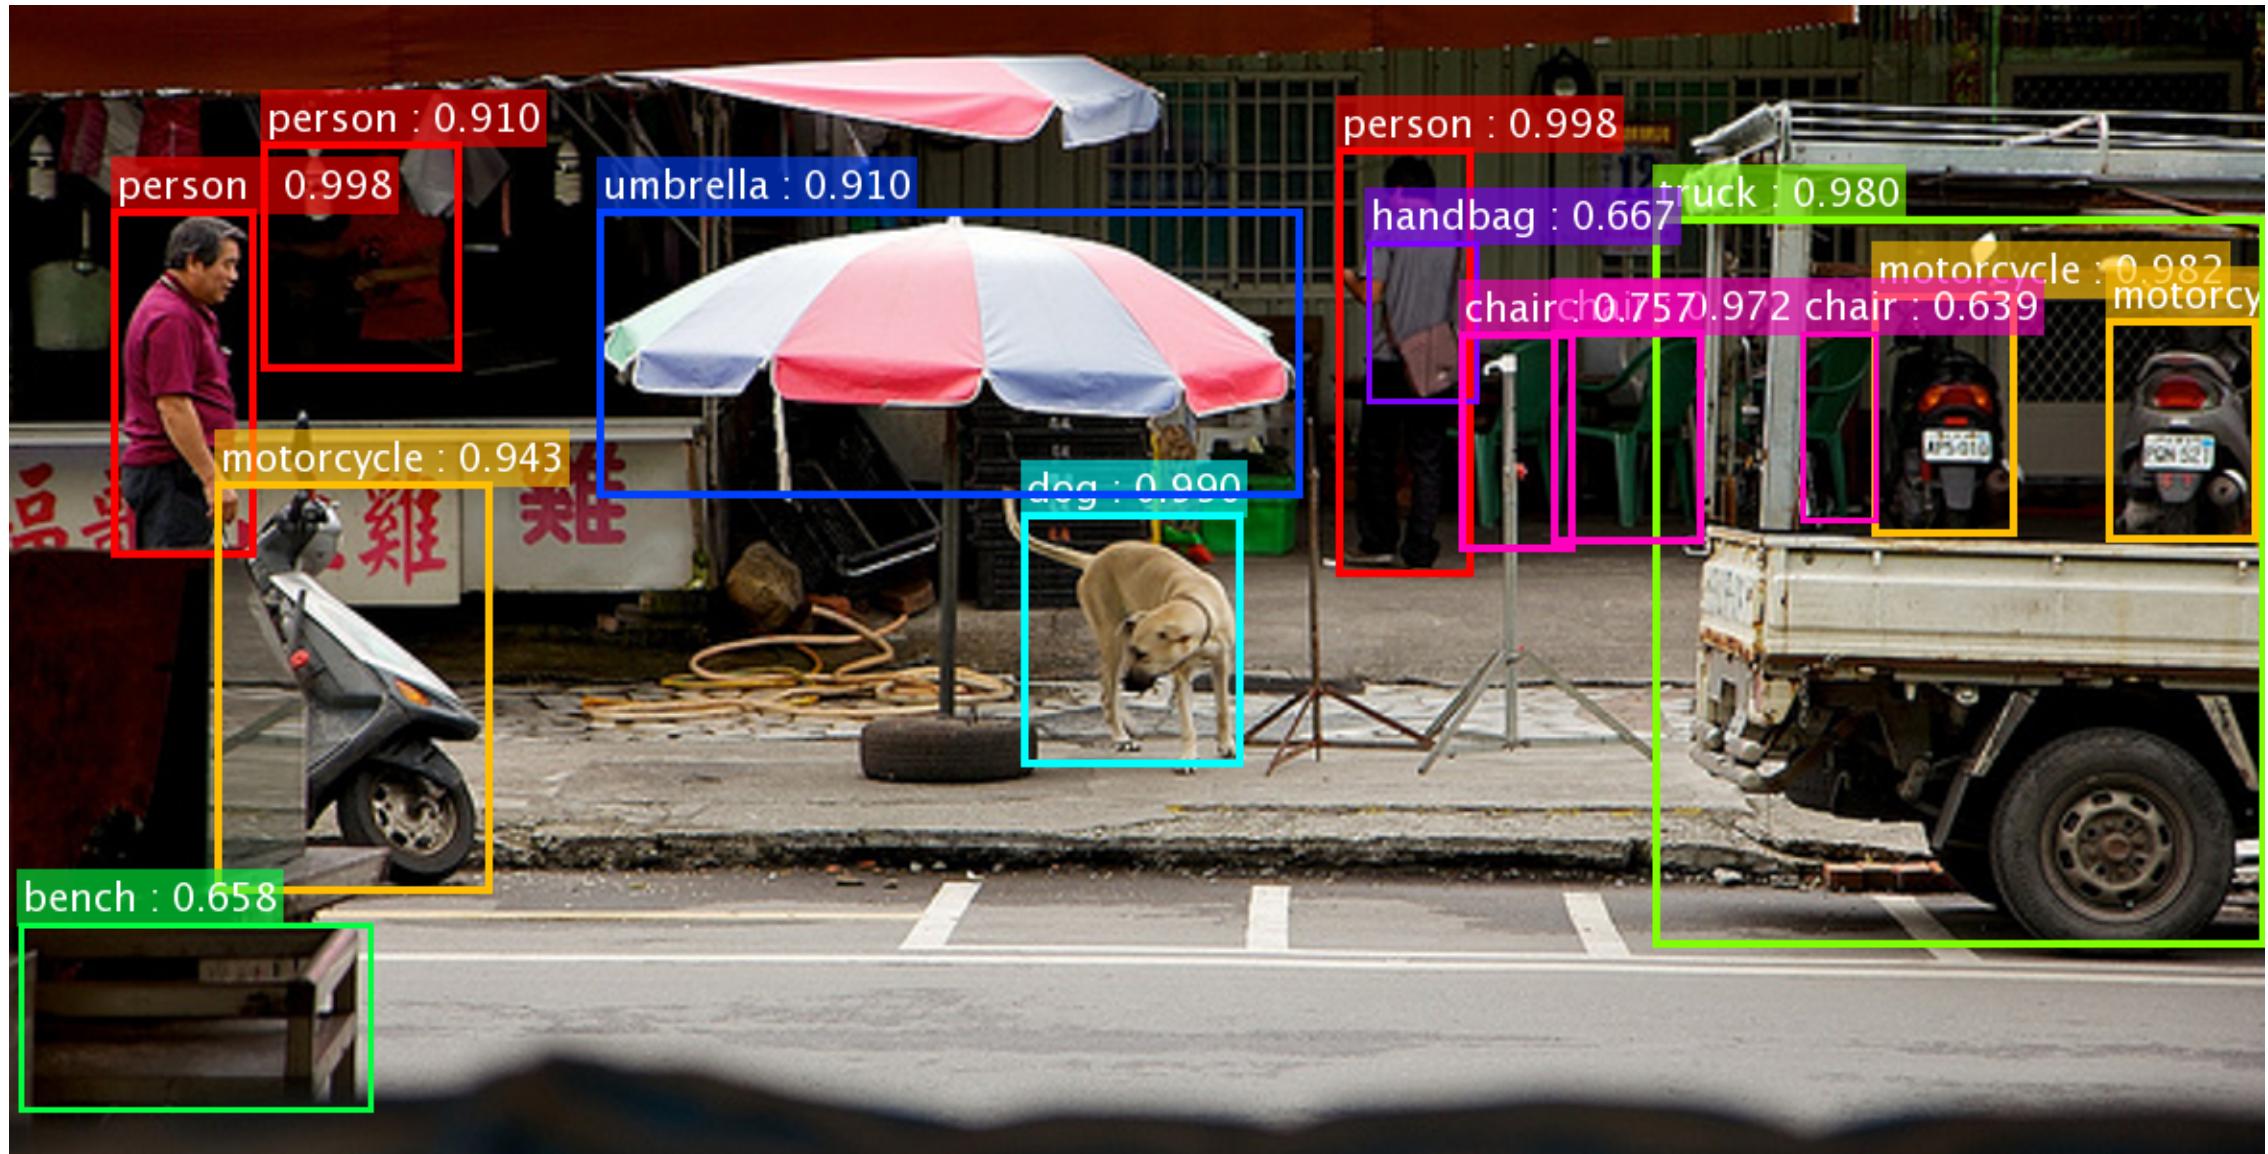
\*the original image is from the COCO dataset

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.  
Shaoqing Ren, Kaiming He, Ross Girshick, & Jian Sun. "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks". NIPS 2015.



\*the original image is from the COCO dataset

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.  
Shaoqing Ren, Kaiming He, Ross Girshick, & Jian Sun. "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks". NIPS 2015.



\*the original image is from the COCO dataset

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.

Shaoqing Ren, Kaiming He, Ross Girshick, & Jian Sun. "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks". NIPS 2015.



Results on real video. Models trained on MS COCO (80 categories).  
(frame-by-frame; no temporal processing)

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.  
Shaoqing Ren, Kaiming He, Ross Girshick, & Jian Sun. "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks". NIPS 2015.

# More Visual Recognition Tasks

ResNet-based methods lead on these benchmarks (incomplete list):

- ImageNet classification, detection, localization
- MS COCO detection, segmentation
- PASCAL VOC detection, segmentation
- Visual Question Answering Challenge 2016
- Human pose estimation [Newell et al 2016]
- Depth estimation [Laina et al 2016]
- Segment proposal [Pinheiro et al 2016]
- ...

	mean	aero	bicycle	bird	boat	bottle	bus	car	cat	chair	motorcycle	sofa	train	tv
DeepLabv2-CRF [?]	79.7	92.6	60.4	91.6	63.4	76.3	95.0	88.4	92.1	84.2	88.3	95.1	92.2	92.1
CASIA_SegResNet_CRF_COCO [?]	79.3	93.8	74.1	91.6	65.5	89.5	95.1	88.3	92.1	84.2	88.3	95.1	92.2	92.1
Adelaide_VeryDeep_FCN_VOC [?]	79.1	91.9	48.1	93.4	69.3	75.5	94.2	87.5	92.1	84.2	88.3	95.1	92.2	92.1
LRR_4x_COCO [?]	78.7	93.2	44.2	89.4	65.4	74.9	95.3	87.0	92.1	84.2	88.3	95.1	92.2	92.1
CASIA_IVA_OASeg [?]	78.3	93.8	41.9	89.4	67.5	71.5	94.6	85.3	89.1	84.2	88.3	95.1	92.2	92.1
Oxford_TVGV_HO_CRF [?]	77.9	92.5	59.1	90.3	70.6	74.4	92.4	84.1	88.1	84.2	88.3	95.1	92.2	92.1
Adelaide_Context_CNN_CRF_COCO [?]	77.8	92.9	39.6	84.0	67.9	75.3	92.7	83.8	89.1	84.2	88.3	95.1	92.2	92.1

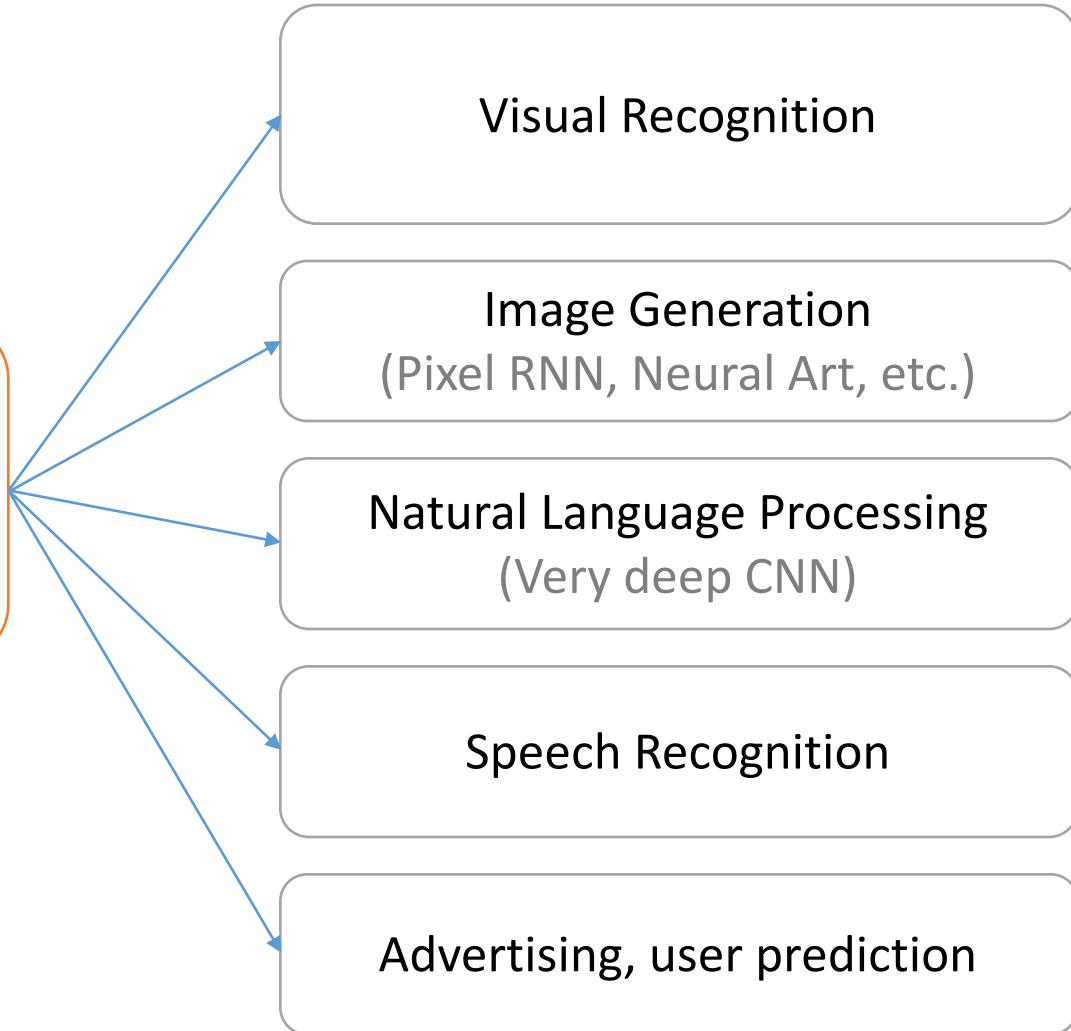
**PASCAL segmentation leaderboard**

	mean	aero	bicycle	bird	boat	bottle	bus	car	cat	chair	motorcycle	sofa	train	tv
Faster RCNN, ResNet (VOC+COCO) [?]	83.8	92.1	88.4	84.0	75.9	71.4	86.3	87.8	94.2	84.2	88.3	95.1	92.2	92.1
R-FCN, ResNet (VOC+COCO) [?]	82.0	89.5	88.3	83.7	75.8	70.7	85.3	86.3	91.1	84.2	88.3	95.1	92.2	92.1
OHEN+FRCN, VGG16, VOC+COCO [?]	80.1	90.1	87.7	79.5	65.0	60.5	80.1	85.0	92.3	84.2	88.3	95.1	92.2	92.1
SSD500 VGG16 VOC + COCO [?]	78.7	89.1	85.7	78.9	63.3	57.0	85.3	84.1	92.3	84.2	88.3	95.1	92.2	92.1
HFM_VGG16 [?]	77.5	88.8	85.1	76.8	64.8	61.4	85.0	84.1	90.0	84.2	88.3	95.1	92.2	92.1
IFRN_07+12 [?]	76.6	87.8	83.9	79.0	64.5	58.9	82.2	82.0	91.4	84.2	88.3	95.1	92.2	92.1
ION [?]	76.4	87.5	84.7	76.8	63.8	58.3	82.6	79.0	90.9	84.2	88.3	95.1	92.2	92.1

**PASCAL detection leaderboard**

# More Applications

ResNets have shown outstanding or promising results on:



# Resources

- Models and Code
  - <https://github.com/KaimingHe/deep-residual-networks>
- Many available implementations  
(see <https://github.com/KaimingHe/deep-residual-networks>)
  - Facebook AI Research's Torch ResNet:  
<https://github.com/facebook/fb.resnet.torch>
    - Torch, CIFAR-10, with ResNet-20 to ResNet-110, training code, and curves: code
    - Lasagne, CIFAR-10, with ResNet-32 and ResNet-56 and training code: code
    - Neon, CIFAR-10, with pre-trained ResNet-32 to ResNet-110 models, training code, and curves: code
    - Torch, MNIST, 100 layers: blog, code
    - A winning entry in Kaggle's right whale recognition challenge: blog, code
    - Neon, Place2 (mini), 40 layers: blog, code
    - .....

# References

## Classification

- “ImageNet Classification with Deep Convolutional Neural Networks”, Krizhevsky et al. NIPS 2012
- “Visualizing and Understanding Convolutional Networks”, Zeiler & Fergus. ECCV 2014
- “Very Deep Convolutional Networks for Large-Scale Image Recognition”, Simonyan & Zisserman. ICLR 2015
- “Going deeper with convolutions”, Szegedy et al. CVPR 2015
- “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”, Szegedy et al. ICML 2015
- “Deep Residual Learning for Image Recognition”, He et al. CVPR 2016

## Detection

- “Rich feature hierarchies for accurate object detection and semantic segmentation”, Girshick et al. CVPR 2014
- “Spatial Pyramid Pooling in Deep Convolutional Networks for Visual Recognition”, He et al. ECCV 2014
- “Fast R-CNN”, Girshick. ICCV 2015
- “Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks”, Ren et al. NIPS 2015

## Segmentation

- “Fully Convolutional Networks for Semantic Segmentation”, Long et al. CVPR 2015
- “Learning to Segment Object Candidates”, Pinhero et al. NIPS 2015

## Language

- “Long-term Recurrent Convolutional Networks for Visual Recognition and Description”, Donahue et al. CVPR 2015
- “Deep visual-semantic alignments for generating image descriptions”, Karpathy & Fei-Fei. CVPR 2015

## Super-Resolution

- “Learning a Deep Convolutional Network for Image Super-Resolution”, Dong et al. ECCV 2014

## Neural Art

- “A Neural Algorithm of Artistic Style”, Gatys et al. CVPR 2016

## Generative models

- “Generative Adversarial Nets”, Goodfellow et al. NIPS 2015